

# Lecture 11: Transforming Conceptual Schema to Logical Schema

---

CSX3006 DATABASE SYSTEMS

ITX3006 DATABASE MANAGEMENT SYSTEMS

# Outline

---

- Why multiple database design steps?
- Transformation of ER model to relational schema
  - Representing Strong Entity Set
  - Representing Composite Attributes
  - Representing Multi-valued attributes
  - Representing Weak Entity Sets
  - Representing 1:1 Relationship Sets
  - Representing 1:M Relationship Set
  - Representing M:M Relationship Sets
  - Relationship Representation and Enforcement of Constraints
  - Representing N-ary Relationships Sets
  - Representing Unary Relationship Sets
  - Representing Inheritance Hierarchy
  - Representing Aggregation

# Conceptual Schema vs. Logical Schema

---

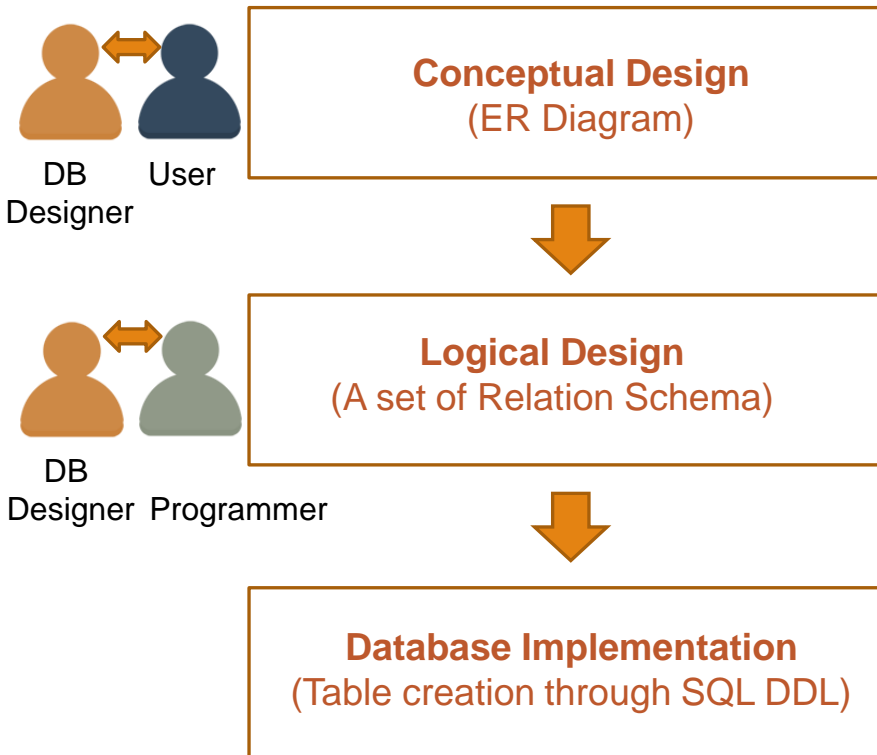
- Conceptual Schema
  - High level abstractions and constructs
    - More natural expressions of the data requirements and business rules
    - Captures the *semantics of data* in an enterprise being modelled
    - Captures some of the *business rules* in the enterprise
  - ER model is mostly adapted for designing the conceptual schema

# Conceptual Schema vs. Logical Schema

---

- Logical Schema
  - ‘Middle-level’ abstractions and constructs
    - Captures the *data requirements of applications* in a format that is *independent of underlying physical file structures, data structure and organizations*.
  - **Relational model** is mostly adapted for designing the logical schema

# Why multiple design steps?



- No need to worry about performance aspects or limitations of a specific DBMS;
  - Focus is on producing a correct model of the real world
- No DBMS can directly support the high level conceptual schema **efficiently**
  - Logical schema are more appropriate for *expressing application logic* by application programmers.
- Note: the conceptual design is not invalidated if a different DBMS is used later


# So what's required?

---

- Transform the conceptual schema into the logical schema supported by a chosen DBMS.

## ER Model

- Entities and Relationships
- Key Constraints, Mapping Cardinality, Participation Constraints



## Relational Model

- relations (tables)
- Primary Key, Candidate Keys, Foreign Keys (Referential Integrity), Unique, Not Null

# Transformation of ER model to relational schema

---

- How to represent Strong Entity Sets in relational schema
- How to represent Composite Attributes in relational schema
- How to represent Multi-valued Attributes in relational schema
- How to represent Weak Entity Sets in relational schema
- How to represent Relationship Sets in relational schema
  - Effects of Cardinality Mappings in the transformation process
  - Effect of Participation Constraints in the transformation process
  - Transformation of ER level Constraints into Relational level constraints
- Representing Generalization/Specialization Hierarchy
- Representing Aggregation

# Representing Strong Entity Set - 1


---

- A **Strong Entity Set** is represented as a **relation** (a **table**) at logical level
  1. **Create a table** with the name of the **strong entity set** being transformed
  2. Every attribute is transformed into a **column** of the table
  3. The **key attribute** (or a set of attributes in the case of composite key) is made **primary key** of the table

| ER Model          |   | Relational Model |
|-------------------|---|------------------|
| Strong entity set | ➡ | Table            |
| Simple attribute  | ➡ | Column           |
| Key attribute     | ➡ | Primary key      |



# Representing Strong Entity Set - 2



ER Model

Employee = ( employee id, name, address, phone)

Relational Model

```
create table employee (  
    employee_id          char(10),  
    name                 char(25),  
    address              char(30),  
    phone                char(30),  
    primary key (employee_id)  
);
```

SQL DDL


# Representing Composite Attributes - 1

---

- **Composite attributes** are flattened out by creating a *separate attributes for each* component attribute

| ER Model             |   | Relational Model     |
|----------------------|---|----------------------|
| Composite attributes | ➡ | Separated attributes |

# Representing Composite Attributes - 2



**Employee = ( employee\_id, first\_name, initial, last\_name, address, phone)**


```
create table employee (  
    employee_id          char(10),  
    first_name           char(15),  
    initial               char(1),  
    last_name            char(15),  
    address               char(30),  
    phone                 char(10),  
    primary key (employee_id)  
);
```

# Representing Multi-valued attributes - 1

- A multi-valued attribute  $M$  of entity set  $E$  is represented by a separate relation  $EM$ 
  - Schema  $EM$  has the following attributes:
    - Attributes: the PK of  $E$ , multi-valued attribute  $M$
    - Primary key (PK)
      - The PK of  $E$  + the multi-valued attribute  $M$
    - Foreign key (FK)
      - Reference to the PK of the  $E$ .
      - This FK CANNOT BE NULL

| ER Model   | Relational Model   |
|--|--|
| Multi-valued attribute $M$<br>Of entity set $E$<br>-- e.g., $E = (\underline{K}, M, N, O)$<br>K is PK. | 1. Revised relation $E$ , and<br>2. A separate relation $EM$ , with FK<br>-- e.g., $E_{\text{revised}} = (\underline{K}, N, O)$<br>$EM = (\underline{K}, \underline{M})$ |

# Representing Multi-valued attributes - 2



Employee = ( employee id, name, address)

Employee\_phone = ( employee id, number)

```
create table employee (  
  employee_id      char(10),  
  name             char(15),  
  address          char(30),  
  primary key (employee_id)  
);
```

```
create table employee_phone (  
  employee_id      char(10),  
  number           char(10),  
  primary key (employee_id, number),  
  foreign key (employee_id) references  
    employee(employee_id)  
  on delete cascade  
  on update cascade
```

# Representing Multi-valued attributes - 3

Employee = ( employee\_id, name, address)

Employee\_phone = ( employee\_id, number)

| employee_id    | name            | address         |
|----------------|-----------------|-----------------|
| 101            | John            | 2/51            |
| <del>102</del> | <del>Sara</del> | <del>35/8</del> |
| 103            | Pete            | 1/20            |

| employee_id | number     |
|-------------|------------|
| 101         | 0811111111 |
| 101         | 0822222222 |
| 102         | 0833333333 |

- What happens when an employee is **deleted** from the Employee table?

```
create table employee (  
  employee_id      char(10),  
  name             char(15),  
  address          char(30),  
  primary key (employee_id)  
);
```

```
create table employee_phone (  
  employee_id      char(10),  
  number           char(10),  
  primary key (employee_id, number),  
  foreign key (employee_id) references  
    employee(employee_id)  
    on delete cascade  
    on update cascade  
);
```

# Representing Multi-valued attributes - 4

Employee = ( employee\_id, name, address)

Employee\_phone = ( employee\_id, number)

| employee_id | name | address |
|-------------|------|---------|
| 101         | John | 2/51    |
| 102         | Sara | 35/8    |
| 103         | Pete | 1/20    |

- What happens when an employee is **deleted** from the Employee table?
  - All corresponding tuples from the employee\_phone must be **deleted**

| employee_id | number     |
|-------------|------------|
| 101         | 0811111111 |
| 101         | 0822222222 |
| 102         | 0833333333 |

```
create table employee (  
  employee_id      char(10),  
  name             char(15),  
  address          char(30),  
  primary key (employee_id)  
);
```

```
create table employee_phone (  
  employee_id      char(10),  
  number           char(10),  
  primary key (employee_id, number),  
  foreign key (employee_id) references  
    employee(employee_id)  
  on delete cascade  
  on update cascade
```

# Representing Multi-valued attributes - 5

Employee = ( employee\_id, name, address)

Employee\_phone = ( employee\_id, number)

- What happens when an employee's id of the Employee table is **updated**?

| employee_id | name | address |
|-------------|------|---------|
| 101         | John | 2/51    |
| 102         | Sara | 35/8    |
| 103         | Pete | 1/20    |

| employee_id | number     |
|-------------|------------|
| 101         | 0811111111 |
| 101         | 0822222222 |
| 102         | 0833333333 |

```
create table employee (  
  employee_id      char(10),  
  name             char(15),  
  address          char(30),  
  primary key (employee_id)  
);
```

```
create table employee_phone (  
  employee_id      char(10),  
  number           char(10),  
  primary key (employee_id, number),  
  foreign key (employee_id) references  
    employee(employee_id)  
  on delete cascade  
  on update cascade  
);
```



# Representing Multi-valued attributes - 6

**Employee = ( employee\_id, name, address)**

**Employee\_phone = ( employee\_id, number)**

| employee_id | name | address |
|-------------|------|---------|
| 101         | John | 2/51    |
| 102         | Sara | 35/8    |
| 103         | Pete | 1/20    |

| employee_id | number     |
|-------------|------------|
| 101         | 0811111111 |
| 101         | 0822222222 |
| 102         | 0833333333 |

- What happens when an employee's id of the Employee table is updated?
  - All corresponding tuples from the Employee\_phone must be updated

```
create table employee (  
  employee_id      char(10),  
  name            char(15),  
  address         char(30),  
  primary key (employee_id)  
);
```

```
create table employee_phone (  
  employee_id      char(10),  
  number          char(10),  
  primary key (employee_id, number),  
  foreign key (employee_id) references  
    employee(employee_id)  
    on delete cascade  
    on update cascade  
);
```

# Representing Multi-valued attributes - 7

Employee = ( employee\_id, name, address)

Employee\_phone = ( employee\_id, number)

| employee_id | name | address |
|-------------|------|---------|
| 101         | John | 2/51    |
| 102         | Sara | 35/8    |
| 103         | Pete | 1/20    |

| employee_id | number     |
|-------------|------------|
| 101         | 0811111111 |
| 101         | 0822222222 |
| 102         | 0833333333 |
| 109         | 0899999999 |

- Can we insert a tuple to Employee\_phone WITHOUT having a corresponding Employee?

```
create table employee (  
  employee_id      char(10),  
  name             char(15),  
  address          char(30),  
  primary key (employee_id)  
);
```

```
create table employee_phone (  
  employee_id      char(10),  
  number           char(10),  
  primary key (employee_id, number),  
  foreign key (employee_id) references  
    employee(employee_id)  
    on delete cascade  
    on update cascade  
);
```

# Representing Multi-valued attributes - 8

Employee = ( employee\_id, name, address)

Employee\_phone = ( employee\_id, number)

FK: NOT NULL

- Can we insert a tuple to Employee\_phone WITHOUT having a corresponding Employee?
  - Cannot. (This rule can be enforced by making the foreign key to be NOT NULL)

Implies UNIQUE NOT NULL, but  
'number' alone needs not be unique

```
create table employee (
    employee_id    char(10),
    name           char(15),
    address        char(30),
    primary key (employee_id)
);
```


| employee_id | name | address |
|-------------|------|---------|
| 101         | John | 2/51    |
| 102         | Sara | 35/8    |
| 103         | Pete | 1/20    |

| employee_id | number     |
|-------------|------------|
| 101         | 0811111111 |
| 101         | 0822222222 |
| 102         | 0833333333 |
| 109         | 0899999999 |

Cannot!!

```
create table employee_phone (
    employee_id    char(10),
    number         char(10),
    primary key (employee_id, number),
    foreign key (employee_id) references
        employee(employee_id)
        on delete cascade
        on update cascade
);
```

# Representing Multi-valued Composite attributes -1



- What should be the **primary key** of the **employee\_phone** table?

Employee = ( employee\_id, name, address )

Employee\_phone = ( *employee\_id*, number, type, location )

## Representing Multi-valued Composite attributes -2

---

- Depends on the business logic!
  - **Case 1:** an employee can have *multiple phone numbers*, but a number is *guaranteed to be unique* within the scope of an employee.
    - **PK = (employee\_id, number)**
  - **Case 2:** an employee may have *two phones in two different locations*, but they happen to have the *same number*!
    - **PK = (employee\_id, number, location)**


• **Think of Referential Integrity Constraints Settings Required!**

# Representing Weak Entity Sets - 1

- A weak entity set is made into a single relation (table) and the identifying relationship set is captured as a foreign key *referring the primary key of owner set*.
- Two tables linked by a **foreign key** reflecting the **identifying set** ← weak entity set
- **The foreign key cannot be null**
  - When the (strong) identifying entity is deleted (or updated), all weak entities depending on the strong entity must be deleted (or updated)

| ER Model  | Relational Model  |
|---|---|
| <p>O: owner entity set,<br/>W: weak entity set with<br/>identifying relationship set I</p> <p><math>O = (\underline{A}, B)</math><br/><math>I = (A, X)</math><br/><math>W = (X, Y)</math></p> | <p>A relation W with FK from the owner entity set</p> <p><math>O = (\underline{A}, B)</math><br/><math>W = (\underline{A}, X, Y)</math></p> |

# Representing Weak Entity Sets - 2



Loan = ( loan\_number, amount)

Payment = ( loan\_number, pay\_number, pay\_amount, pay\_datetime)

FK: NOT NULL

# Representing Weak Entity Sets - 3

Loan = ( loan\_number, amount)

Payment = ( loan\_number, pay\_number, pay\_amount, pay\_datetime)

FK: NOT NULL

```
create table loan (  
  loan_number    char(10),  
  amount        numeric(9,2),  
  primary key (loan_number)  
);
```

```
create table payment (  
  loan_number    char(10),  
  pay_number     numeric(4),  
  pay_amount     numeric(9,2),  
  pay_datetime   timestamp,  
  primary key (loan_number, pay_number),  
  foreign key (loan_number) references  
    loan(loan_number)  
    on delete cascade  
    on update cascade  
);
```

- **loan\_number can't be null** since it's part of the primary key
- **Implicitly NOT NULL**



# Remark: constraints of FK

- We must consider if we can specify the min. and max. cardinalities on each FK.
  - Min. cardinality = 1 → specify by using 'NOT NULL' constraint
  - Max. cardinality = 1 → specify by using 'UNIQUE' constraint
  - Min. cardinality = 0 → no need to specify
  - Max. cardinality = n → no need to specify
    - However, for Max. cardinality = 1, if the FK is also PK of the table, it automatically implies 'UNIQUE'.

| Contract_ID | Job        | Start_Date  | End_Date    | Staff_ID |
|-------------|------------|-------------|-------------|----------|
| 100         | Programmer | 1 Jan. 2019 | 2 Jan. 2020 | 201901   |
| 101         | IT Support | 1 Jan. 2020 | 2 Jan. 2021 | 202001   |
| 102         | Programmer | 1 Feb. 2020 | 1 Feb. 2021 | 202002   |


**FK**

{ Min. cardinality = ?  
Max. cardinality = ?

# Representing 1:1 Relationship Sets – 1

## One-to-One Relationship Sets


---



- An entity on the left can be associated with at most one entity on the right *and vice versa*.

# Representing 1:1 Relationship Sets – 2

## One-to-One Relationship Sets



Alternative 1

$E12 = (\underline{pk1}, a1, \underline{pk2}, a2)$

Alternative 2

$E1 = (\underline{pk1}, a1)$

$E2 = (\underline{pk2}, a2, \underline{fk1})$

Alternative 3

$E1 = (\underline{pk1}, a1, \underline{fk2})$

$E2 = (\underline{pk2}, a2)$

Alternative 4

$E1 = (\underline{pk1}, a1)$


$R = (\underline{fk1}, \underline{fk2})$

$E2 = (\underline{pk2}, a2)$

- Total participation on both sides ( 1..1 ; 1..1 )
- Total-Partial; Partial-Total ( 1..1 ; 0..1 ) ( 0..1 ; 1..1 )
- Partial on both sides ( 0..1 ; 0..1 )

# Representing 1:1 Relationship Sets – 1

## Total Participation at BOTH sides




- It is SOMETIMES possible to **subsume** one entity set into the other
  - The choice *depends on which is the most important entity set* (more attributes, better key, semantic nature of them).

$$E12 = (\underline{pk1}, a1, \overline{pk2}, a2)$$

- The PK of the new combined entity => the *original more important entity set***
- The candidate key => the *key of the subsumed entity set***
- If there are any attributes in common, the duplicates are removed.

# Representing 1:1 Relationship Sets – 2

## Total Participation at BOTH sides



- **Example:** One-to-One between member and mobilePhone BOTH total participation
  - Each member is *required* to have *one and only one* mobile number.
  - Each mobile number is *required* to belong to *one and only one* member.


**member = ( member\_id, name, address, mobile\_number, mobile\_operator )**

# Representing 1:1 Relationship Sets – 3

## Total Participation at BOTH sides

---

- So here is another design alternative even at ER level
  - Pros and cons of each alternatives?
  - WHEN should NOT Combine?



`member = ( member_id, name, address, mobile_number, mobile_operator )`


# Representing 1:1 Relationship Sets – 4

## Total Participation at BOTH sides

---


- **When not to combine**

- the two entity types *represent different entities* in the 'real world'.
- the entities *participate in* very different relationships *with other entities*.
- efficiency considerations:
  - fast responses are required
  - different patterns of updating occur to the two different entity types.
    - *E.g., Yearly contract staff* (sign a new contract in yearly basis)



# Representing 1:1 Relationship Sets – 5

## Total Participation at BOTH sides



- **If NOT COMBINED** into a single entity set,
  - The PK of one entity set becomes the FK in the other.
  - **Note: should NOT have a FK in BOTH entity.**
  - **Referential Integrity on the FK**
    - **must NOT be NULL** → enforces the **MIN**imum cardinality of 1
    - **must be UNIQUE** → enforces the **MAX**imum cardinality of 1

E1 = (pk1, a1)

E2 = (pk2, a2, **fk1**)

OR

E1 = (pk1, a1, **fk2**)


E2 = (pk2, a2)



# Representing 1:1 Relationship Sets – 6

## Total Participation at BOTH sides

---



**Choice 1 (combine entity set)**

**Staff = ( emp\_id, name, cont\_number, position, salary, start, end)**

# Side Note: Integrity Constraints - 1

## Choice 2

Staff = ( emp\_id, name, **cont\_number** )

Referential Integrity on the FK

**NOT NULL, UNIQUE**

ON UPDATE CASCADE, ON DELETE RESTRICT

Contract = ( cont\_number, position, salary, start, end )

- **CAN**
  - Insert a new contract *without* a corresponding staff member
    - Need application logic to make sure that there will be a staff member which will be associated with this new contract.
  - Delete a staff member; consequently, the associated contract will *not* have the corresponding staff member anymore
    - Need application logic to ensure that the contract is deleted
- **CANNOT**
  - Insert a staff member which does not have a contract
  - Insert a staff member having the same contract number as another existing staff member (must be one staff per contract).

## Side Note: Integrity Constraints - 2

### Choice 3

Staff = ( emp\_id, name)

Contract = ( cont\_number, position, salary, start, end, *emp\_id* )

Referential Integrity on the FK


NOT NULL, UNIQUE

ON UPDATE CASCADE, ON DELETE CASCADE

- **CAN**
  - Insert a staff member *without* a contract
  - Delete a contract; consequently, a staff member will become contract-less!
- **CANNOT**
  - Insert a contract which is not associated with a staff member
  - Insert a contract associated with the same staff as an existing contract
    - Each staff is associated with maximum of one contract

# Representing 1:1 Relationship Sets – 1

## One-to-One with one Partial participation



- Two entities sets are kept in *separate tables*, and
- The *association* between them are made through a FK from the *total side to the partial side*
  - Avoid null value in FK.


NOT NULL

|                                      |   |
|--------------------------------------|---|
| E1 = ( <u>pk1</u> , a1, <i>fk2</i> ) | T |
| E2 = ( <u>pk2</u> , a2)              | P |

A red arrow points from the *fk2* attribute in the E1 tuple to the pk2 attribute in the E2 tuple, illustrating the foreign key relationship from the total side (E1) to the partial side (E2).

# Representing 1:1 Relationship Sets – 2

## One-to-One with one Partial participation



Staff = ( emp\_id, name)

Referential Integrity on the FK

**NOT NULL, UNIQUE**

ON UPDATE CASCADE, **ON DELETE CASCADE**

Note: what to do  
on delete and  
on update  
depends on application!

Contract = ( cont\_number, position, salary, start, end, emp\_id )

| Cont_number | position | salary | start | end | Emp_id (FK) |
|-------------|----------|--------|-------|-----|-------------|
| C111        | ...      | ...    | ...   | ... | 101         |
| c222        | ...      | ...    | ...   | ... | 102         |
| c333        | ...      | ....   | ...   | ... | 105         |

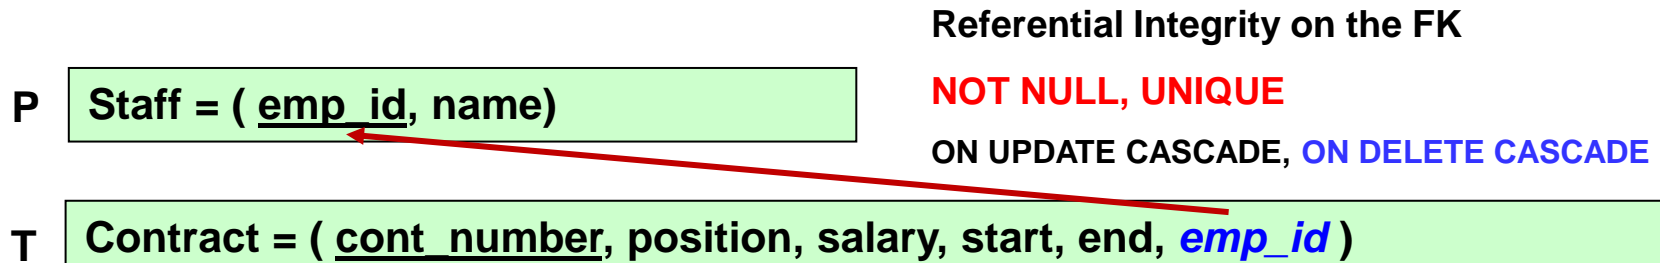
WONG

| Emp_id | name | Cont_number (FK) |
|--------|------|------------------|
| 101    | A    | C111             |
| 102    | B    | c222             |
| 103    | C    | null             |

**Not good!**

# Side Note: Integrity Constraints


---



- **NOT NULL on the FK** -> Cannot insert a contract without a staff member
- **UNIQUE on the FK** -> A staff member can be associated with only one contract
- **ON UPDATE (DELETE) CASCADE** -> Updating (deleting) an emp\_id in Staff table will result in updating (deleting) also in the Contract table.
- What if a staff member is inserted without a contract? ok
- What if a contract is deleted? What happens to the Staff member?
  - No need for any special action → since Staff is in the partial participation

# Representing 1:1 Relationship Sets - 1


## One-to-One; Partial participation on BOTH sides



- **Choice 1: (2 Tables)**
  - Two entities sets are kept in *separate tables*, and
  - The *association* between them are *made through a FK*
  - The PK of one entity set becomes the FK in the other *but not* vice versa.
    - The FK must be Allowed to be NULL (why? Partial participation!)
- **Choice 2: (3 Tables)**
  - **Create a relation** (table) for the relationship, and
  - associate the entities through the FKs in the newly created table
    - What should be the primary key of the third table?

# Representing 1:1 Relationship Sets – 2

## One-to-One; Partial participation on BOTH sides



### Choice 1a:

employee = ( emp\_id, name, **laptop\_serial** )

laptop = ( laptop\_serial, brand, model )

Referential Integrity on the foreign key  
**UNIQUE**, but NULL must be allowed  
ON UPDATE CASCADE, **ON DELETE SET NULL**

### Choice 1b:

employee = ( emp\_id, name )


laptop = ( laptop\_serial, brand, model, **emp\_id** )

Referential Integrity on the foreign key  
**UNIQUE**, but NULL must be allowed  
ON UPDATE CASCADE, **ON DELETE SET NULL**



# Representing 1:1 Relationship Sets – 3

## One-to-One; Partial participation on BOTH sides



automatically implied  
'UNIQUE NOT NULL'  
from being PK of the  
table.

Choice 2:

employee = ( emp\_id, name)

emp\_laptop = ( emp\_id, laptop\_serial )

laptop = ( laptop\_serial, brand, model)

Referential Integrity on each of the FKs

**UNIQUE** and **NOT NULL**

**ON UPDATE CASCADE, ON DELETE CASCADE**


(Note: 'UNIQUE' implies max. cardinality = 1,  
'NOT NULL' implies min. cardinality = 1)

```
create table emp_laptop (
    emp_id char(10)
    laptop_serial char(15) unique not null
    primary key (emp_id),
    foreign key (emp_id) references
        employee(emp_id)
        on delete cascade
        on update cascade,
    foreign key (laptop_serial) references
        laptop(laptop_serial)
        on delete cascade
        on update cascade
);
```

# Representing 1:1 Relationship Sets - 1

## Descriptive Attributes

- Reposition descriptive attributes to **either** entity set
  - But place them with the **table having the FK**
- Example 1:




Staff = ( emp\_id, name )

Contract = ( cont number, position, salary, start, end, **emp\_id**, **signed\_on** )

# Representing 1:1 Relationship Sets - 2

## Descriptive Attributes

- Example 2



employee = ( emp\_id, name)


emp\_laptop = ( emp\_id, laptop\_serial, duration)

laptop = ( laptop\_serial, brand, model)

# Representing 1:M Relationship Set

## One-to-Many Relationship Sets


---



- An entity on the left can be associated with **many** entities on the right
- An entity on the right can be associated with **at most one** entity on the left
- The methods of representing **One-to-Many** Relationship sets equally apply for representing **Many-to-One**

# Representing 1:M Relationship Set - 1

## Total Participation on BOTH sides




- The two entity sets are represented by *two separate tables*, and
- The *association* are made through a *FK*
  - A FK is created on the MANY side table referencing the PK of the ONE side table
  - Referential Integrity on the FK
    - **must NOT be NULL** → enforces the total participation
      - The minimum cardinality = 1
- The maximum cardinality of 1 is enforced by the fact that PK2 is the **primary key** (*automatically implies 'UNIQUE'*);
  - Each E2 entity instance is associated with only one E1 instance

E1 = (pk1, a1)

E2 = (pk2, a2, *fk1*)

# Representing 1:M Relationship Set - 2

## Total Participation on BOTH sides



department = ( dept\_id, budget, building )

employee = ( emp\_id, name, dept\_id )

Referential Integrity on the FK


**NOT NULL**

ON UPDATE CASCADE, ON DELETE ???????

- ❑ An employee CANNOT be inserted without being associated with a department
  - ❑ Enforced by NOT NULL on the FK
- ❑ An employee CANNOT be associated with two or more departments
  - ❑ Enforced by emp\_id being the PK
- ❑ Does not automatically ensure that a department must have one or more employees
  - ❑ Must be regulated through application logic

# Representing 1:M Relationship Set - 1

## Total on Many side and Partial on One side




- The two entity sets are represented by *two separate tables*, and
- The *association* are made through a *FK*
  - A FK is created on the MANY side table referencing the PK of the ONE side table
  - Referential Integrity on the FK
    - **must NOT be NULL** → enforces the total participation
  - The maximum cardinality of 1 is enforced by the fact that PK2 is the primary key;
  - Each E2 entity instance is associated with only one E1 instance

E1 = (pk1, a1)

E2 = (pk2, a2, *fk1*)

# Representing 1:M Relationship Set – 2

## Total on Many side and Partial on One side



teacher = ( teacher\_id, name, office)

project = ( proj\_id, title, *teacher\_id* )

Referential Integrity on the FK

**NOT NULL**


ON UPDATE CASCADE, ON DELETE ???????

- A project **CANNOT** be inserted **without** being associated with a teacher
  - Enforced by **NOT NULL** on the FK
- A project **CANNOT** be associated with **two or more** teachers
  - Enforced by **proj\_id** being the PK
- The **one side is partial** (Teacher may or may not look after any project)
  - There is **no need to put any other constraints**



# Representing 1:M Relationship Set - 1

## Partial on Many side and Total on One side (choice 1)



**Choice 1**  
**(2 Tables)**

E1 = (pk1, a1)

E2 = (pk2, a2, *fk1*)


**Not a good choice!!**

### Referential Integrity on the FK

- **NULL Allowed** – since the many side is partial participation; there may be entities in E2 which are not associated with any entities in E1
- **UNIQUE NOT Possible** – different E2s can be associated with the same E1
- maximum cardinality of 1 on the E1 is enforced by pk2 being the primary key of E2
- However, **the total participation of the E1 (one side) is not enforced automatically**
  - Must be regulated through application logic

# Representing 1:M Relationship Set - 2

## Partial on Many side and Total on One side (choice 1)



branch = (branch id, name, city)

laptop = (laptop id, brand, model, **branch\_id**)


**Not a good choice!!**

**FK**

- Null allowed
- UNIQUE NOT possible

# Representing 1:M Relationship Set - 3

## Partial on Many side and Total on One side (choice 2)



### Choice 2 (3 Tables)

- Note: the PK of R is the PK of the MANY side
- Referential Integrity on fk2
  - UNIQUE, NOT NULL
    - implied automatically since it's the primary key of R
- Referential Integrity on fk1
  - NOT NULL

E1 = (pk1, a1)

R = (fk1, fk2)


E2 = (pk2, a2)

R:

| fk1 | fk2 |
|-----|-----|
| 1   | A   |
| 1   | B   |
| 1   | C   |

# Representing 1:M Relationship Set - 3

## Partial on Many side and Total on One side (choice 2)



- Branch MUST have one or more laptops
- Laptop MAY be owned by a branch, but not more than one branch

| Branch_id | Laptop_id |
|-----------|-----------|
| A         | L1        |
| A         | L2        |
| B         | L3        |

### Choice 2

branch = (branch\_id, name, city)


branchlaptop = (branch\_id, laptop\_id)

laptop = (laptop\_id, brand, model)

```
create table branchlaptop (  
    laptop_id      char(10)  
    branch_id      char(15) not null  
    primary key (laptop_id),  
    foreign key (branch_id) references  
        branch(branch_id),  
    foreign key (laptop_id) references  
        laptop(laptop_id)  
);
```

# Representing 1:M Relationship Set - 1

## Partial Participation on BOTH sides (Choice 1)



**Not a good choice!!**

| pk2 | Fk1 (Pk1) |
|-----|-----------|
| A   | 1         |
| B   | 1         |
| C   | 2         |
| D   | Null      |
| E   | Null      |
| F   | Null      |

- Choice 1 (2 Tables)

- Referential Integrity on the FK


- **NULL Allowed** – since the many side is partial participation; there may be entities in E2 which are not associated with any entities in E1
- **UNIQUE NOT Possible** – different E2s can be associated with the same E1
- maximum cardinality of 1 on the E1 is enforced by pk2 being the primary key of E2

E1 = (pk1, a1)


E2 = (pk2, a2, *fk1*)

# Representing 1:M Relationship Set - 2

## Partial Participation on BOTH sides (Choice 2)




- Choice 2 (3 Tables)
  - Note: the PK of R is the PK of the MANY side
- Referential Integrity on fk2
  - UNIQUE, NOT NULL
    - implied automatically since it's the primary key of R
- Referential Integrity on fk1
  - NOT NULL



# Representing 1:M Relationship Set - 3

## Partial Participation on BOTH sides (Choice 2)



teacher = (teacher\_id, name, office)

advising = (teacher\_id, student\_id)

student = (student\_id, name, cgpa)

```

create table advising (
    student_id      char(7)
    teacher_id      char(6) not null
    primary key (student_id),
    foreign key (student_id) references
        student(student_id),
    foreign key (teacher_id) references
        teacher(teacher_id)
);
    
```

Teacher\_id = {L1, L2, **L3**, **L4**}

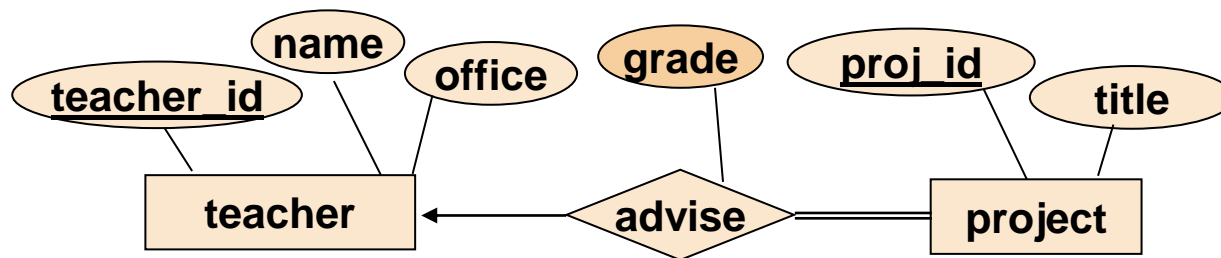
Student\_id = {1,2,3,**4**,**5**}

| Teacher_id | Student_id |
|------------|------------|
| L1         | 1          |
| L2         | 2          |
| L2         | 3          |

# Representing 1:M Relationship Set - 1

## Descriptive Attributes

- Reposition descriptive attributes **to the MANY side**
  - place them with the table having the FK
- Example 1



teacher = ( teacher\_id, name, office )


project = ( proj\_id, title, **teacher\_id**, **grade** )



# Representing 1:M Relationship Set - 2

## Descriptive Attributes

- Example 2




teacher = (teacher\_id, name, office)


advising = (*teacher\_id*, *student\_id*, last\_meeting)

student = (student\_id, name, cgpa)

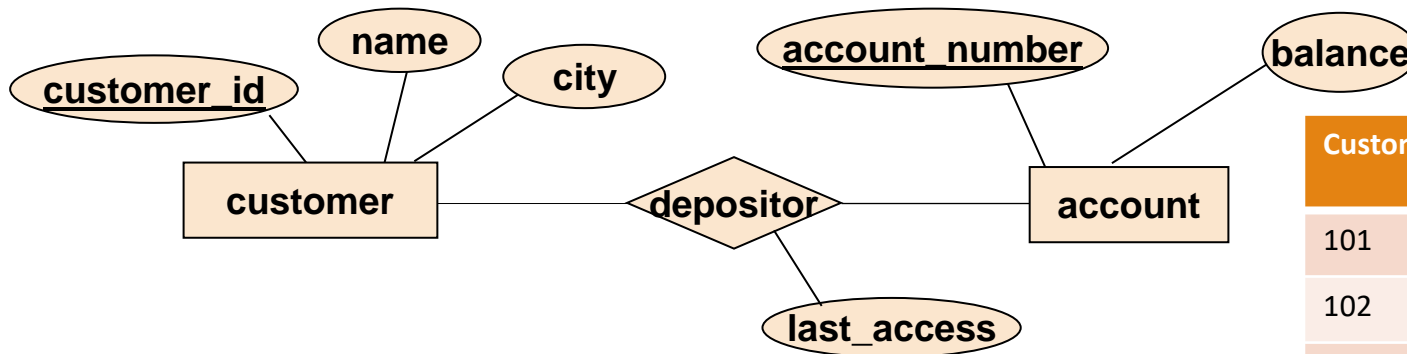
# Representing M:M Relationship Sets - 1



- **Always create a separate table for the relationship**
  - The table will contain the PKs of the entity sets participating in the relationship **plus** any descriptive attributes of the relationship set
  - **The PK of the table will be the UNION of the PKs of the entity sets**
  - Each individual FK **CANNOT be unique** (Many-to-Many)



# Representing M:M Relationship Sets - 2



| Customer_id | Account_number | Last_access |
|-------------|----------------|-------------|
| 101         | A_111          | Yesterday   |
| 102         | A_111          | Yesterday   |
| 101         | A_222          | Today       |
| 102         | A_222          | Today       |

customer = (customer\_id, name, city)

depositor = (customer\_id, account\_number, last\_access)

account = (account\_number, balance)

```
create table depositor (  
    customer_id      char(7),  
    account_number   char(6),  
    last_access       date,  
    primary key (customer_id, account_number),  
    foreign key (customer_id) references  
        customer(customer_id),  
    foreign key (account_number) references  
        account(account_number)  
);
```

# Relationship Representation and Enforcement of Constraints - 1

---

- **One-to-One**
  - Sometimes can be merged into a single table one when Total-Total
  - Two tables with a FK from one table to the other
    - Total-Total; Total-Partial; Partial-Total
  - Three tables with FK to the other two tables
    - Especially for Partial-Partial or when the relationship has attributes
- **One-to-Many (Many-to-One)**
  - Two tables with a FK from MANY side to the ONE side
    - Total-Total; Many (Total) - One (Partial)
  - Three tables with FKs to the other two tables
    - One (Total) – Many (Partial) ; Partial-Partial, and when relationship has attributes

# Relationship Representation and Enforcement of Constraints - 2

---

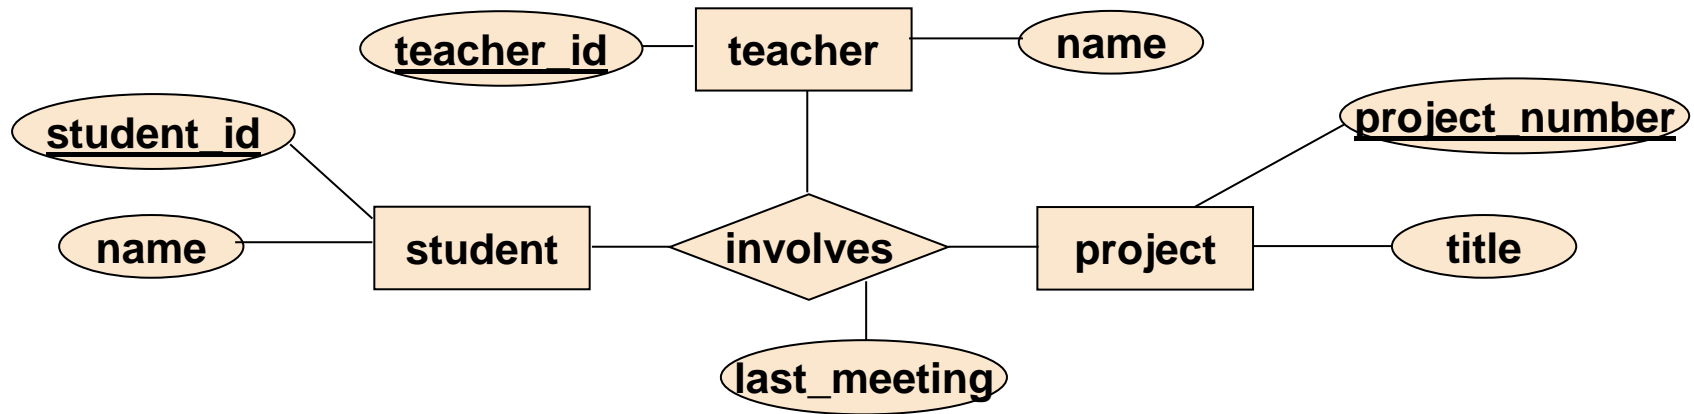
- **Many-to-Many**
  - Always three tables with FKs to the other two tables
- **Realize that NOT all kinds of Cardinality Mapping and Participation Constraints can be automatically enforced**
  - Total Participation (or minimum cardinality of 1) may not be enforceable in certain cases
  - Specific Cardinality mappings like ( 2..8 ) cannot be automatically enforced
    - Need to use other techniques like check, trigger, stored procedures, or application logic

# Representing N-ary Relationships Sets - 1

---

- **Steps of transforming relationship sets having degree  $> 2$** 
  1. Create a table for the relationship
  2. Include the PKs of all the entity sets involved and any other descriptive attributes of the relationship set as the attributes of the table
  3. Create FKs referencing the PK of each entity set participating in the relationship
  4. The PK of the table will be the UNION of the PK of the entity sets involved in the relationship.
    - Except when there is a single arrow
    - The PK of the table will be the UNION of the PKs of entity sets not on the arrow side.
- But first, re-evaluate your ER design to see if you really need the N-ary relationship.

# Representing N-ary Relationships Sets - 2



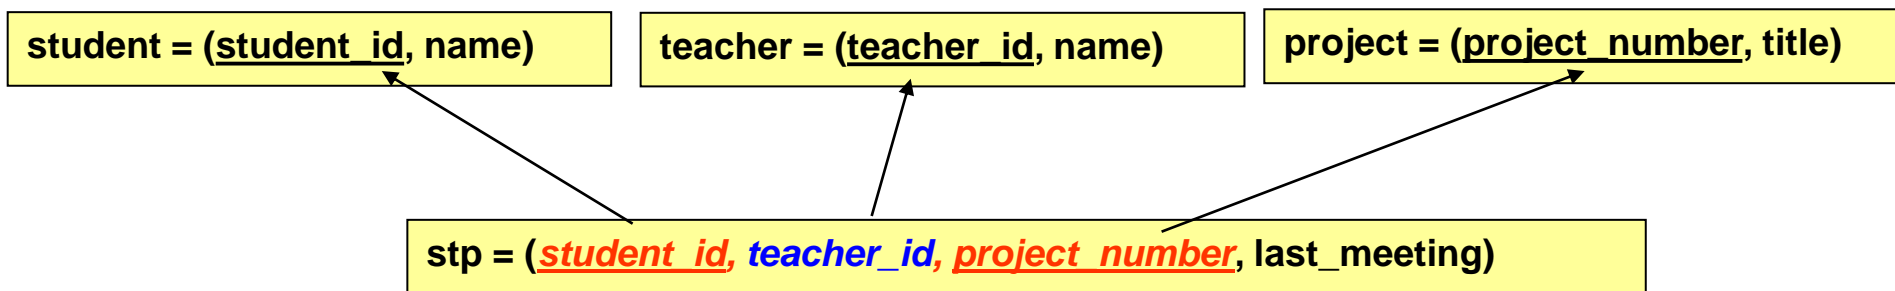
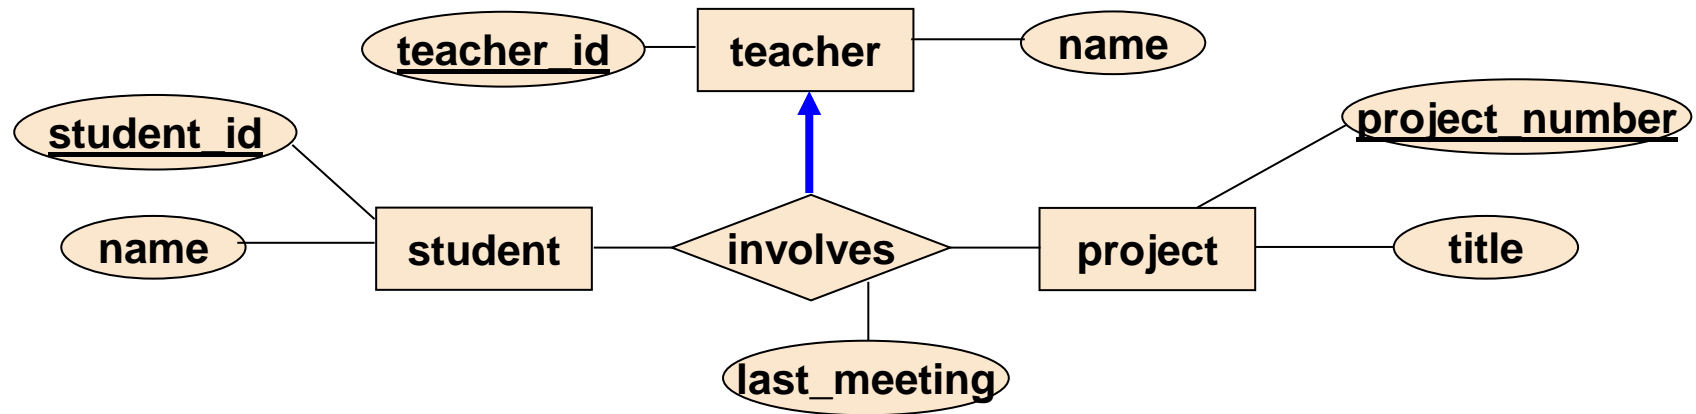
student = (student\_id, name)

teacher = (teacher\_id, name)

project = (project\_number, title)

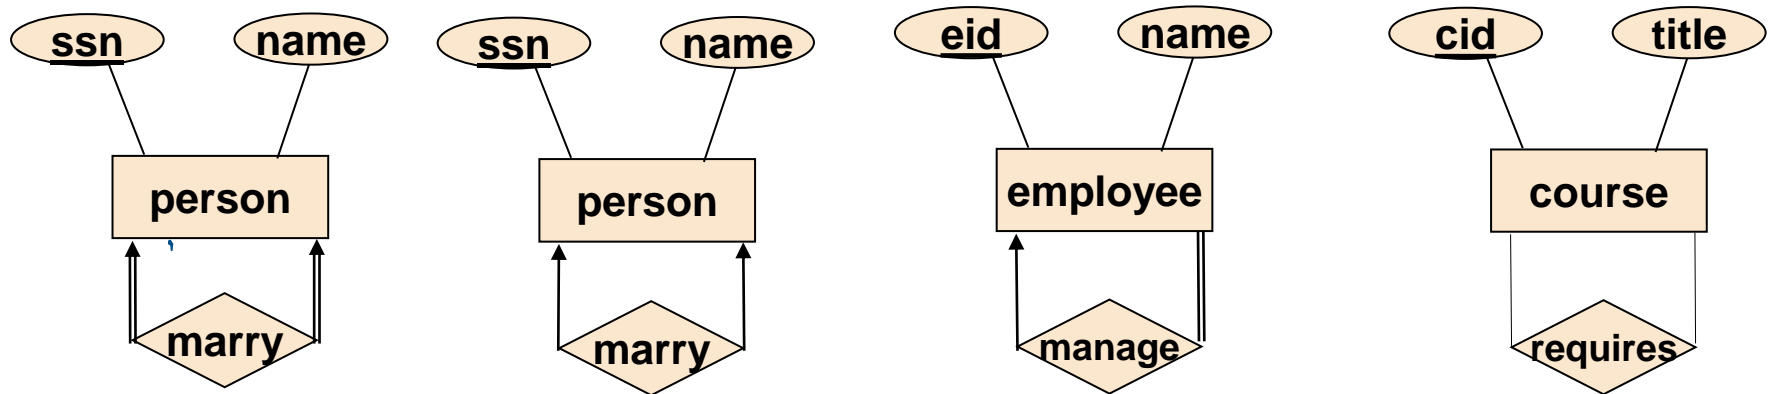
stp = (student\_id, teacher\_id, project\_number, last\_meeting)

# Representing N-ary Relationships Sets - 3





# Representing Unary Relationship Sets - 1



- **Same methods as those of transforming Binary Relationship Sets**
  - Unary Relationship Sets = *binary relationship set, where*
    - *an entity set participates in the same relationship set twice, each with different roles.*

# Transforming Binary Relationship Sets

---

## 1. Determine the Cardinality Mappings

- One-To-One, One-to-Many, and Many-to-One


## 2. Determine the Participation Constraints

- Total vs. Partial (Or Minimum Cardinality)

## 3. Choose the methods for the required choice

- Single table with a FK referring back to its own PK
- Extra table for the relationship having FK to the entity table


# Representing Unary Relationship Sets - 2



person = (ssn, name, spouse)

```
create table person (  
  ssn          char(15),  
  name         char(6),  
  spouse char(15) unique not null,  
  primary key (ssn),  
  foreign key (spouse) references person(ssn)  
);
```

# Representing Unary Relationship Sets - 3



course = (cid, title)

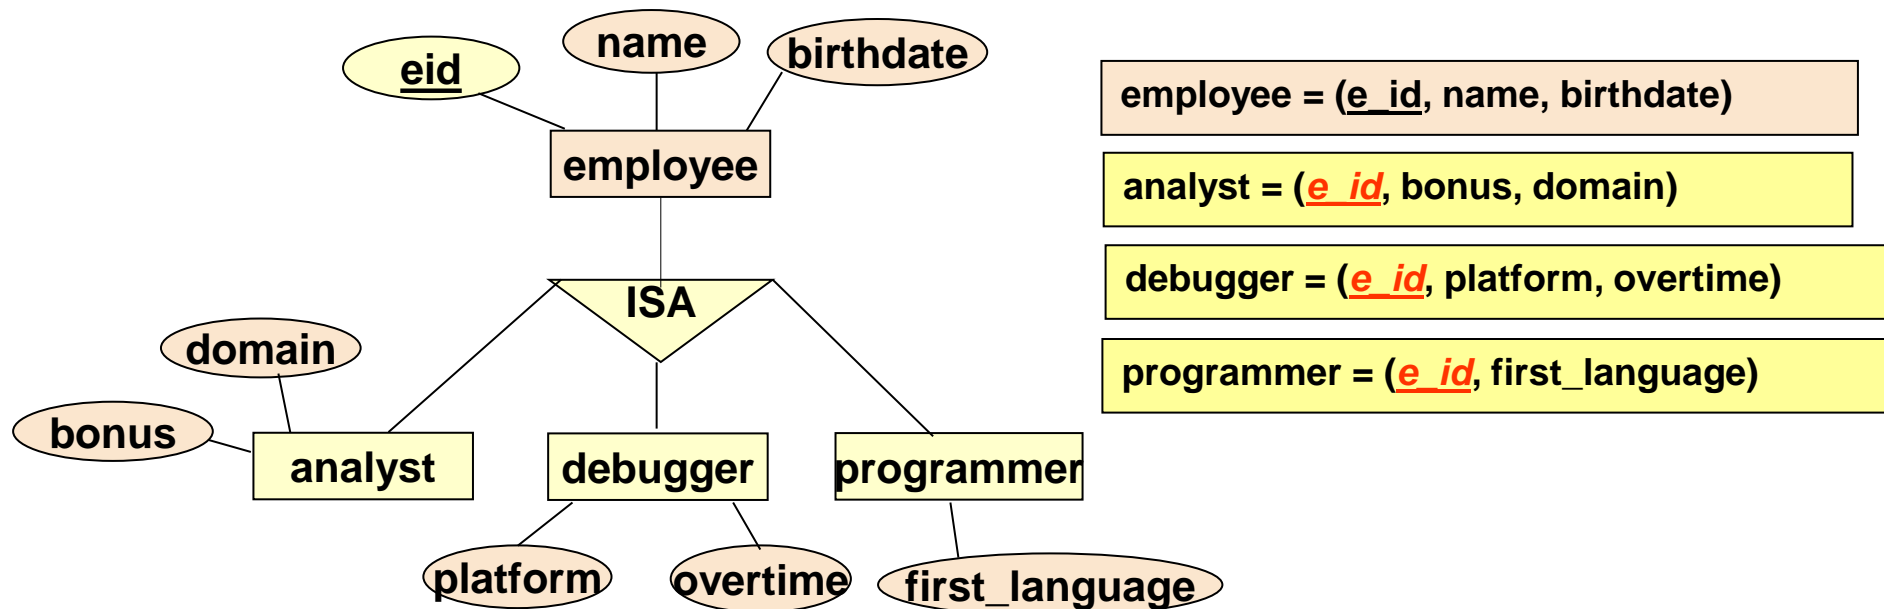
prerequisite = (cid, pre)

```
create table course (  
    cid          char(6),  
    title        char(50),  
    primary key (cid)  
);  
create table prerequisite (  
    cid          char(6),  
    pre          char(6),  
    primary key (cid, pre),  
    foreign key (cid) references course(cid),  
    foreign key (pre) references course(cid)  
);
```

# Representing Inheritance Hierarchy – 1

## Strategy 1 (The most close to the EER)

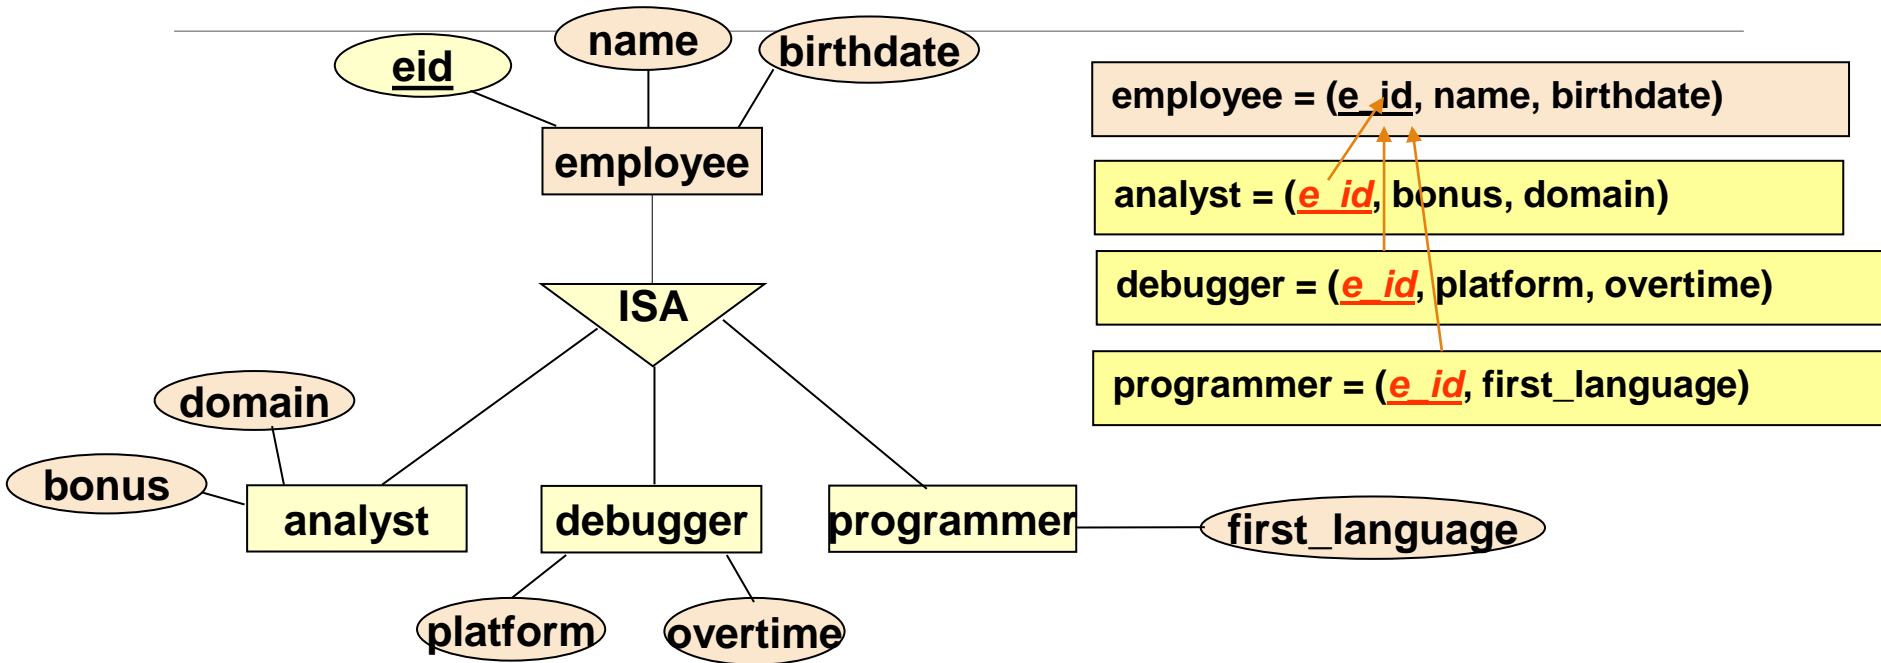
- The PK of the super-class is mapped into each subclass and becomes the PK of the sub-classes.
- The PK of each sub-class is **also made a FK**



□ Note: the enforcement of constraints (overlapping, disjoint, total, partial) need to be done through application logic

# Representing Inheritance Hierarchy – 2

## Strategy 1 (The most close to the EER)

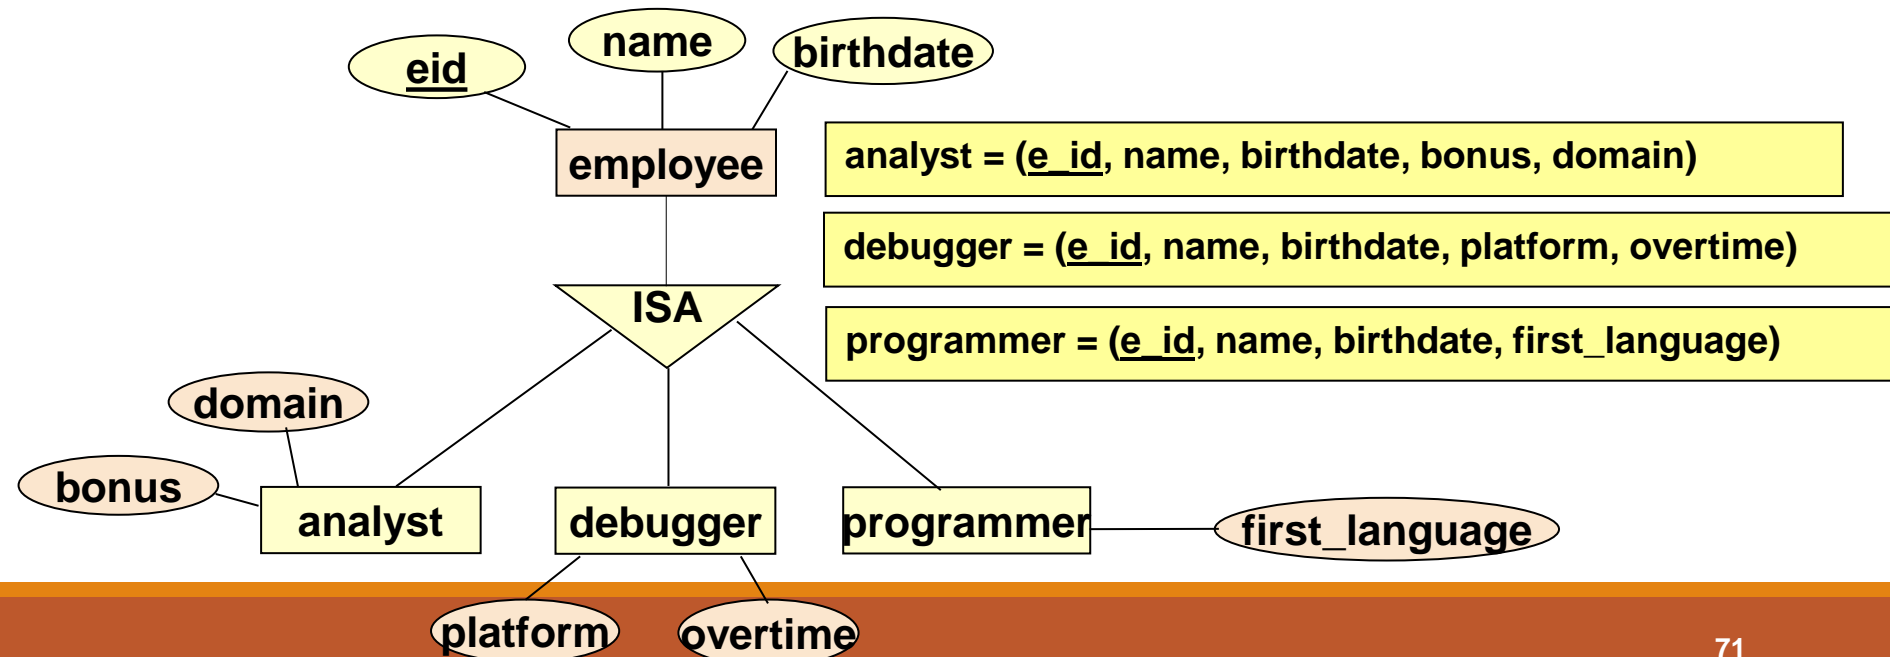


- Referential Integrity Constraints on the FKs
  - **UNIQUE, NOT NULL**
- Enforcement of Total, Partial, Disjoint, Overlapping rules need to be enforced through check, triggers, or stored procedures.
- What should be actions on delete and on update?

# Representing Inheritance Hierarchy – 3

## Strategy 2 (Disregard the hierarchy and represent as separate entity sets)

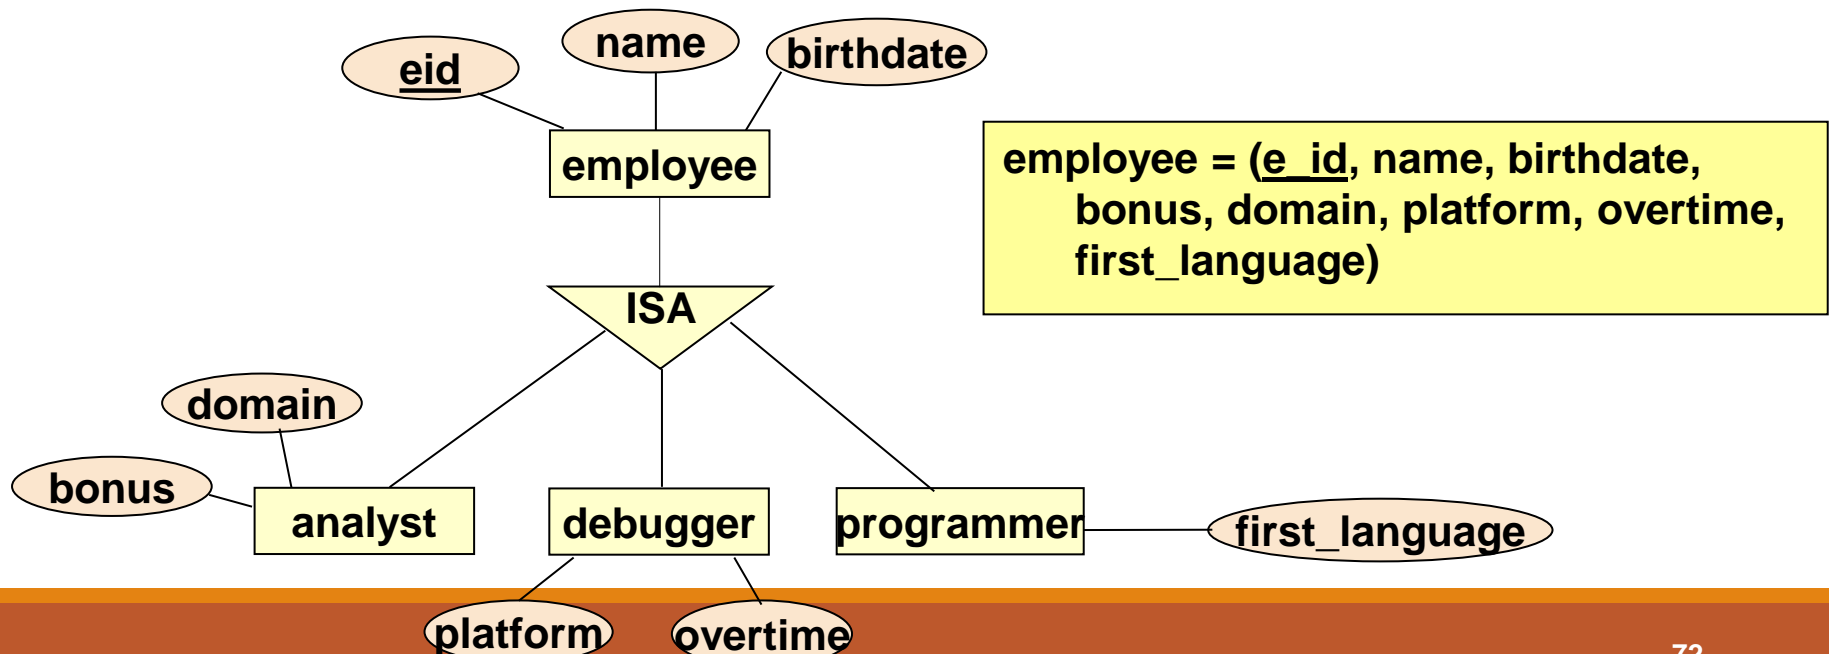
- OK when the *hierarchy is disjoint and total* and when the superclass is not participating any other relationship with other entity sets..
- If overlapping, then result in data redundancy → Data Inconsistency
- Not possible to define a single FK constraints on the super-class if the super-class is participating in other relationships with other entity sets.



# Representing Inheritance Hierarchy – 4

## Strategy 3 (Disregard the hierarchy and represent as a single entity set)

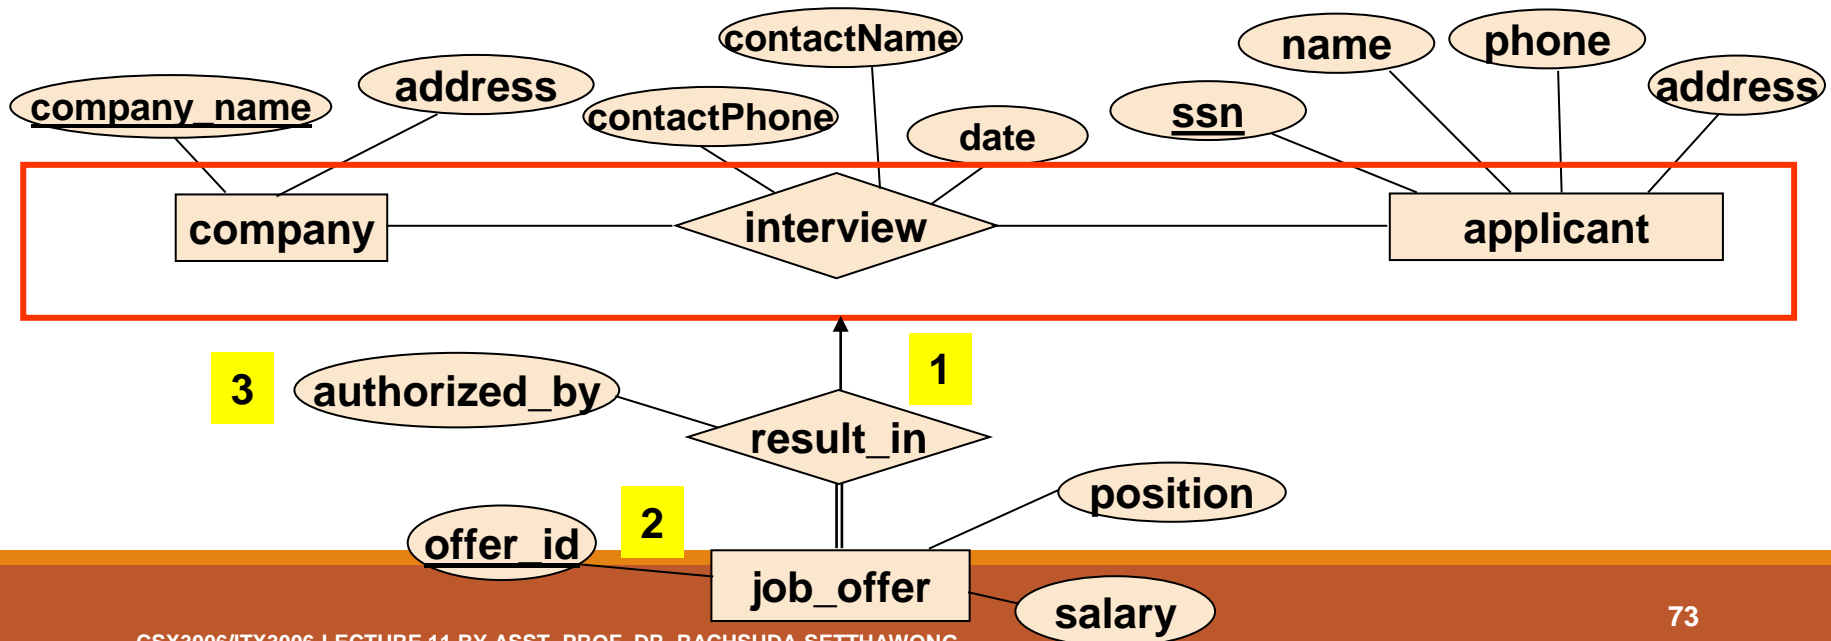
- **Not really a good choice** as there will be many null values
- Furthermore, **Cannot handle situations when the sub-classes are engaged in specific relationships with other entity sets.**
- Only good point is that it avoids joins to get additional information about each subclass
- **Do NOT use this strategy unless there is huge penalty in the performance caused by join operations required by strategy #1**



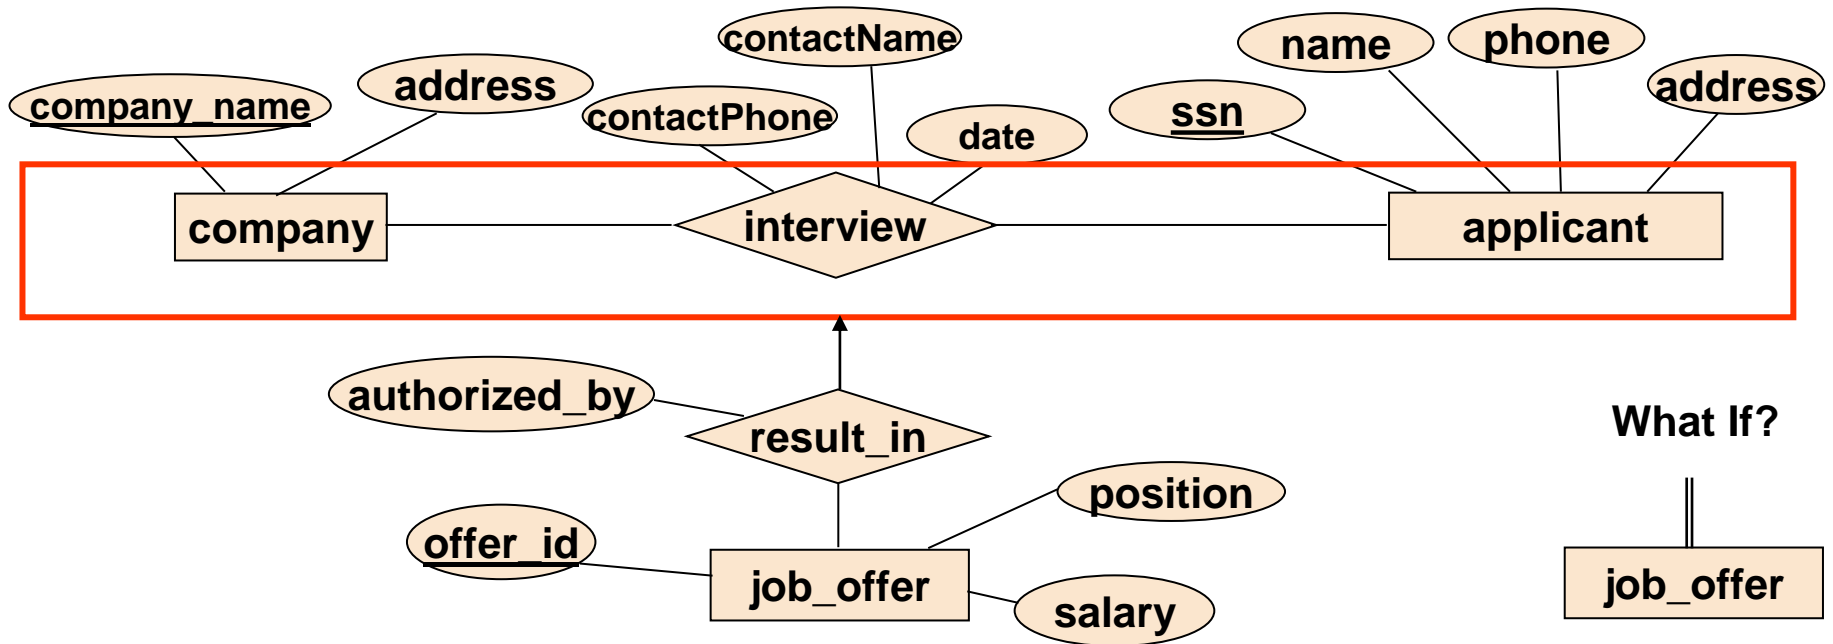


# Representing Aggregation

- Create a table containing
  1. PK of the aggregated relationship,
  2. the PK of the associated entity set
  3. any descriptive attributes
- Note: the **whole aggregation** is represented as a table for the relationship linking the 'aggregated entity set' to the other entity set (e.g. result\_in)



# Representing Aggregation - 2



company = (company\_name, address)

applicant = (ssn, name, phone, address)

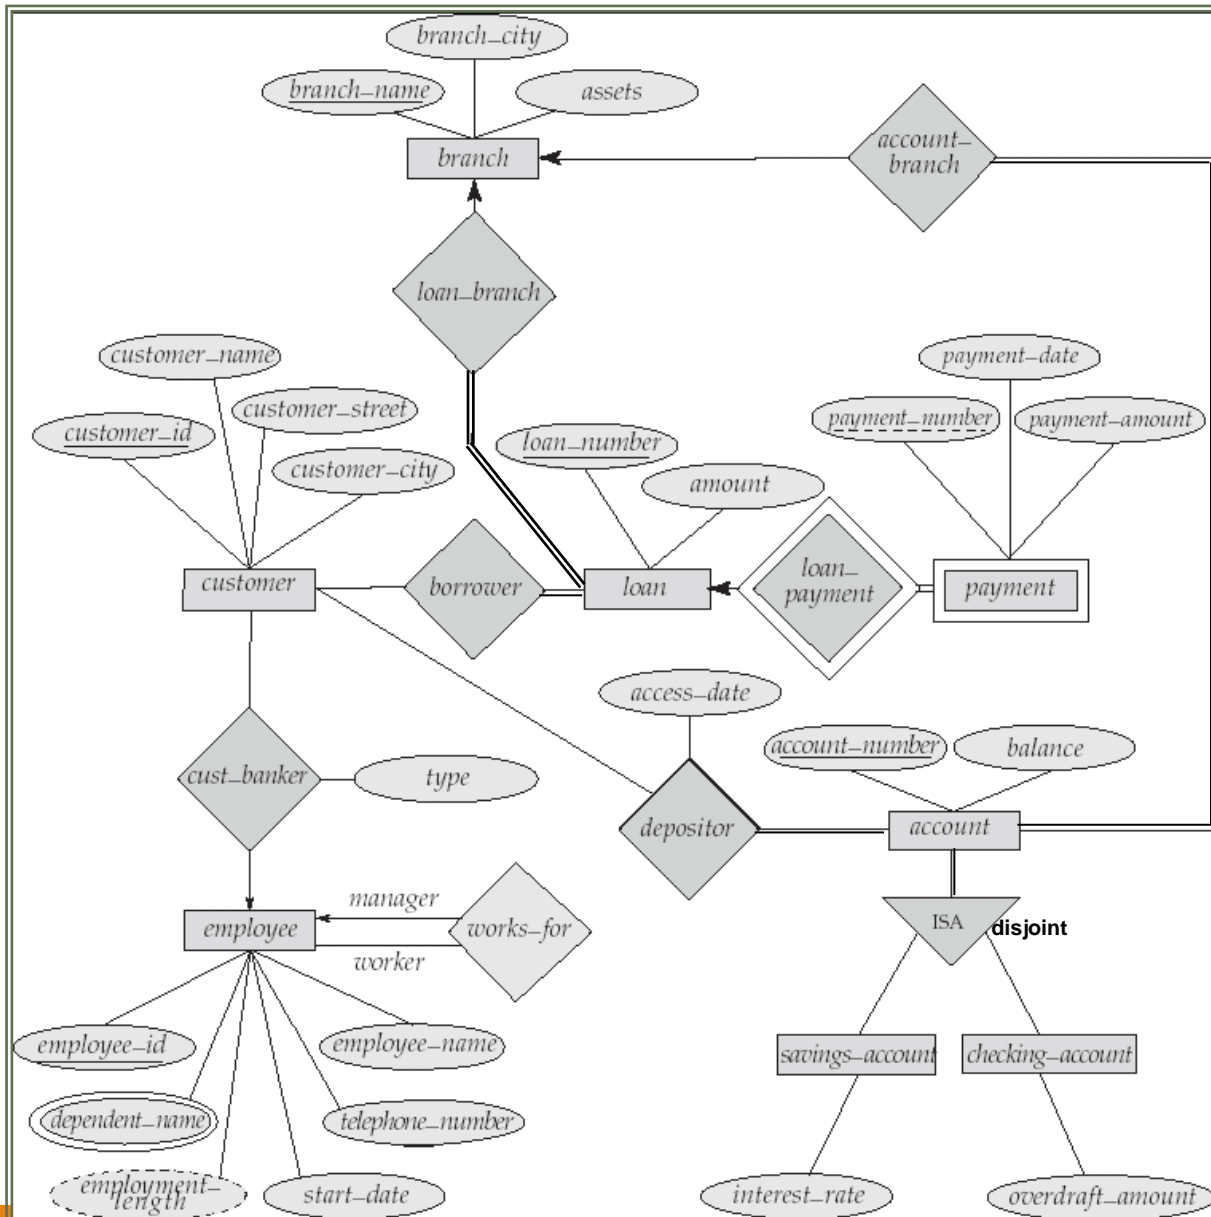
interview = (company\_name, ssn, contactName, contactPhone, date)

result = ( company\_name, ssn, offer\_id, authorized\_by )

job\_offer = (offer\_id, position, salary)

Offer\_id is the PK of result table since it's a Many-to-One relationship from job\_offer to interview

# Self-Exercise #1



- Transform into a set of relational schema
- Define SQL DDL to create the tables
- Carefully consider the integrity constraints specification in DDL.

# Self-Exercise #2

---

- Modify the ER diagram shown in the previous slide to support the banking model where the bank keeps track of deposits and withdrawals from savings and checking accounts.
  - Assume the bank does not care who made the deposits or withdrawals.
  - Assume the bank does not care about methods of deposits or withdrawals, e.g.) Cash Withdrawal, withdrawal from ATM machine 178, online transfer
  - It should be very similar how loan payments are tracked