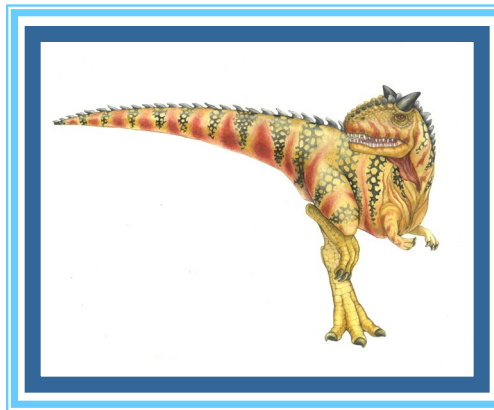


Chapter 3: Processes





Chapter 3: Processes

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication (IPC)
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems





Objectives

- ❑ To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- ❑ To describe the various features of processes, including scheduling, creation and termination, and communication
- ❑ To explore interprocess communication (IPC) using shared memory and message passing
- ❑ To describe communication in client-server systems





Process Concept

- An OS executes a variety of programs:
 - **Batch system** – jobs (assigned work)
 - **Time-shared systems** – user programs or tasks
 - Even on a *single-user system*, a user may be able to run several programs at one time:
 - ▶ a word processor, a Web browser, and an e-mail package.





The Process (1)

- **Process** – a program in execution (its .exe file in main memory, and it is taken by the CPU for execution or waiting for execution); process execution must progress sequentially.
- A **process** has multiple parts:
 - The **program code**, also called the **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing *temporary data* such as
 - ▶ Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time
 - ▶ The structure of a process in memory is shown in **Figure 3.1**.





The Process (2)




Figure 3.1 Process in memory.





The Process (3)

- ❑ It is noticed that a program (.exe file) by itself is not a process.
- ❑ A **program** is a ***passive entity***, such as a file containing a list of instructions stored on a disk memory (often called an **executable file**).
- ❑ In contrast, a **process** is an ***active entity*** in **main memory** with a PC reg. Specifying the next instruction to execute with a set of associated resources.
 - ❑ That is, a program becomes a process only when its executable file is loaded into main memory.





The Process (4)

- Although **two processes** may be associated with the same program, they are considered two separate execution sequences.
 - For instance, several users may be running different copies of the mail program, or the same user invoke many copies of the web browser, are **separate processes**.
 - ▶ and although the text sections are equivalent, the data, heap, and stack sections vary.
- It is also common to have a **process** that **spawns** (generate a new) many processes as it runs.
 - An executable **Java** code is executed within the *Java virtual machine* (JVM). The JVM executes as **a process** that interprets the loaded *Java code* and takes actions.





Process State (1)

- As a **process** executes, it changes state;
 - The state of a process is defined in part by the current activity of that process.
- A **process** may be in one of the following **states**:
 - **New**: The process is being created.
 - **Running**: Instructions are being executed.
 - **Waiting**: The process is waiting for some event to occur (such as an I/O completion or reception of a signal).
 - **Ready**: The process is waiting to be assigned to a processor.
 - **Terminated**: The process has finished execution.





Process State (2)

- Only one process can be running on any processor at any instant. Many processes may be **ready** and **waiting**.
 - The state diagram corresponding to these states is presented in **Figure 3.2**.

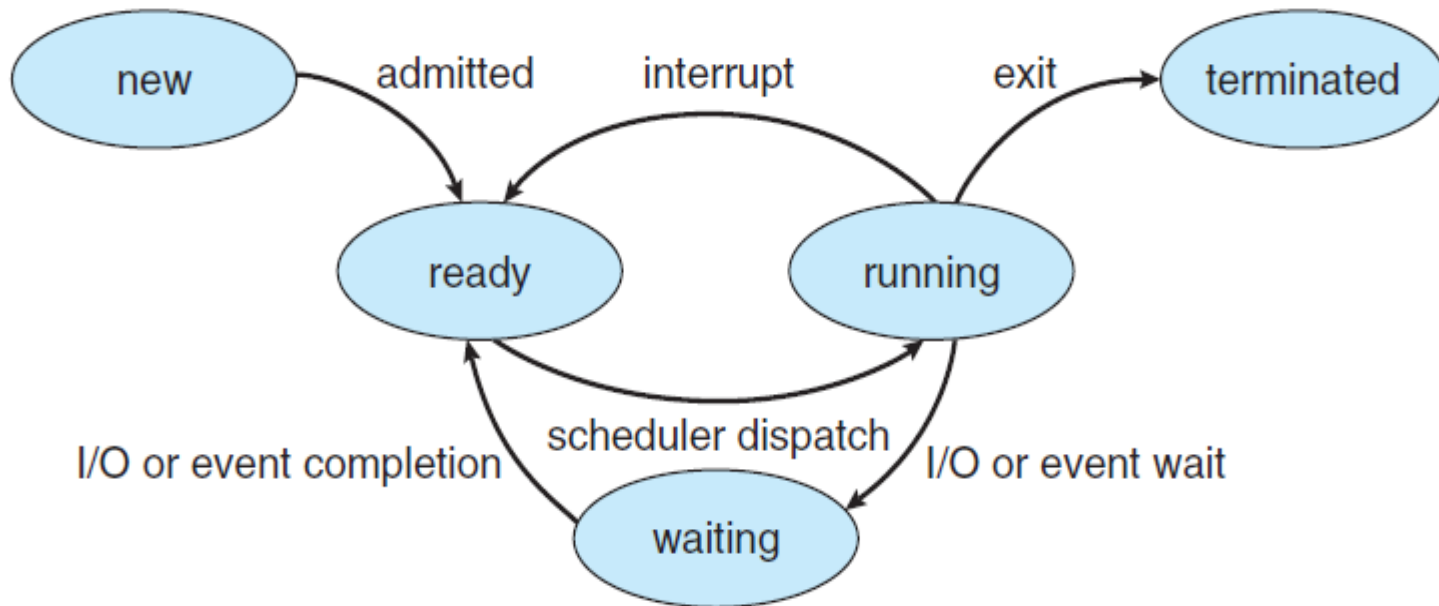


Figure 3.2 Diagram of process state.





Process Control Block (1)

- Each process is represented in the OS by a **process control block (PCB)**—also called a **task control block**.
- A **PCB** is shown in **Figure 3.3**.

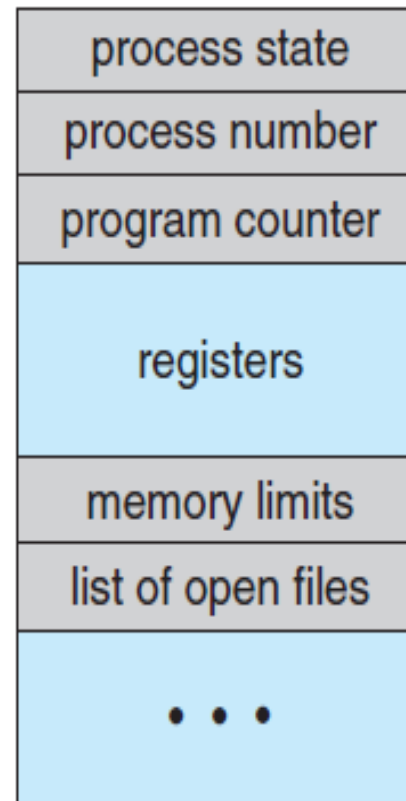


Figure 3.3 Process control block (PCB).





Process Control Block (2)

- A **PCB** contains many pieces of information associated with a specific **process**, including these:
 - **Process state:** The state may be new, ready, running, waiting, halted, and so on.
 - **Program counter:** The counter indicates the address of the next instruction to be executed for this process.
 - **CPU registers:** The size of registers depending on the computer architecture.
 - **CPU-scheduling information:** This information includes a process priority, pointers to scheduling queues, etc.
 - **Memory-management information:** This information includes the value of the *base* and *limit* registers and the *page tables*, etc
 - **Accounting information:** This information includes the amount of CPU and real time used, time limits, account numbers, job or process numbers, etc.
 - **I/O status information:** This information includes the list of I/O devices allocated to the process, a list of open files, and so on.





A CPU Switches from Process to Process




Figure 3.4 Diagram showing CPU switch from process to process.





Threads

- The process model discussed is a **program** that performs a **single thread** of execution.
 - For example, when a process is running a word-processor program, a **single thread of instructions** is being executed at each time.
 - This **single thread of control** allows the process to perform only one task at a time.
- Most modern OSs have extended the process concept to allow a process to have **multiple threads** of execution and thus to perform more than one task at a time.
 - This feature is especially beneficial on **multicore** systems, where **multiple threads can run in parallel**.





Process Scheduling

- The objective of **multiprogramming** is to have some process running at all times, to **maximize CPU utilization**.
- The objective of **time sharing** is to switch the CPU among processes so users can interact with each program while it is running.
- To meet these objectives, the **process scheduler** selects an available process for program execution on the CPU.
- For a **single-processor system**, there will never be more than one running process.
 - If there are more processes, the rest will have to wait until the CPU is free.





Scheduling Queues

- ❑ As **processes** enter the system, they are put into a **job queue**,
 - ❑ which consists of all processes in the system.
- ❑ The processes that are residing in **main memory** and are **ready** and **waiting** to execute are kept on a list called the **ready queue**.
 - ❑ This **queue** is generally stored as a linked list.
 - ❑ A **ready-queue** header contains pointers to the first and final PCBs in the list.
 - ❑ Each PCB includes a pointer field that points to the next PCB in the **ready queue**.
- ❑ The list of processes waiting for a particular **I/O device** is called a **device queue**.





Scheduling Queues

- A common representation of **process scheduling** is a **queuing diagram**, such as that in **Figure 3.6**.
 - Each rectangular box represents a **queue**.
- Two types of queues are present:
 - the **ready queue** and
 - a set of **device queues**
 - The **circles** represent the **resources** that serve the queues, and the **arrows** indicate the **flow of processes** in the system.
- A **new process** is initially put in the **ready queue**.
 - It waits there until it is selected for execution, or **dispatched**.





Scheduling Queues

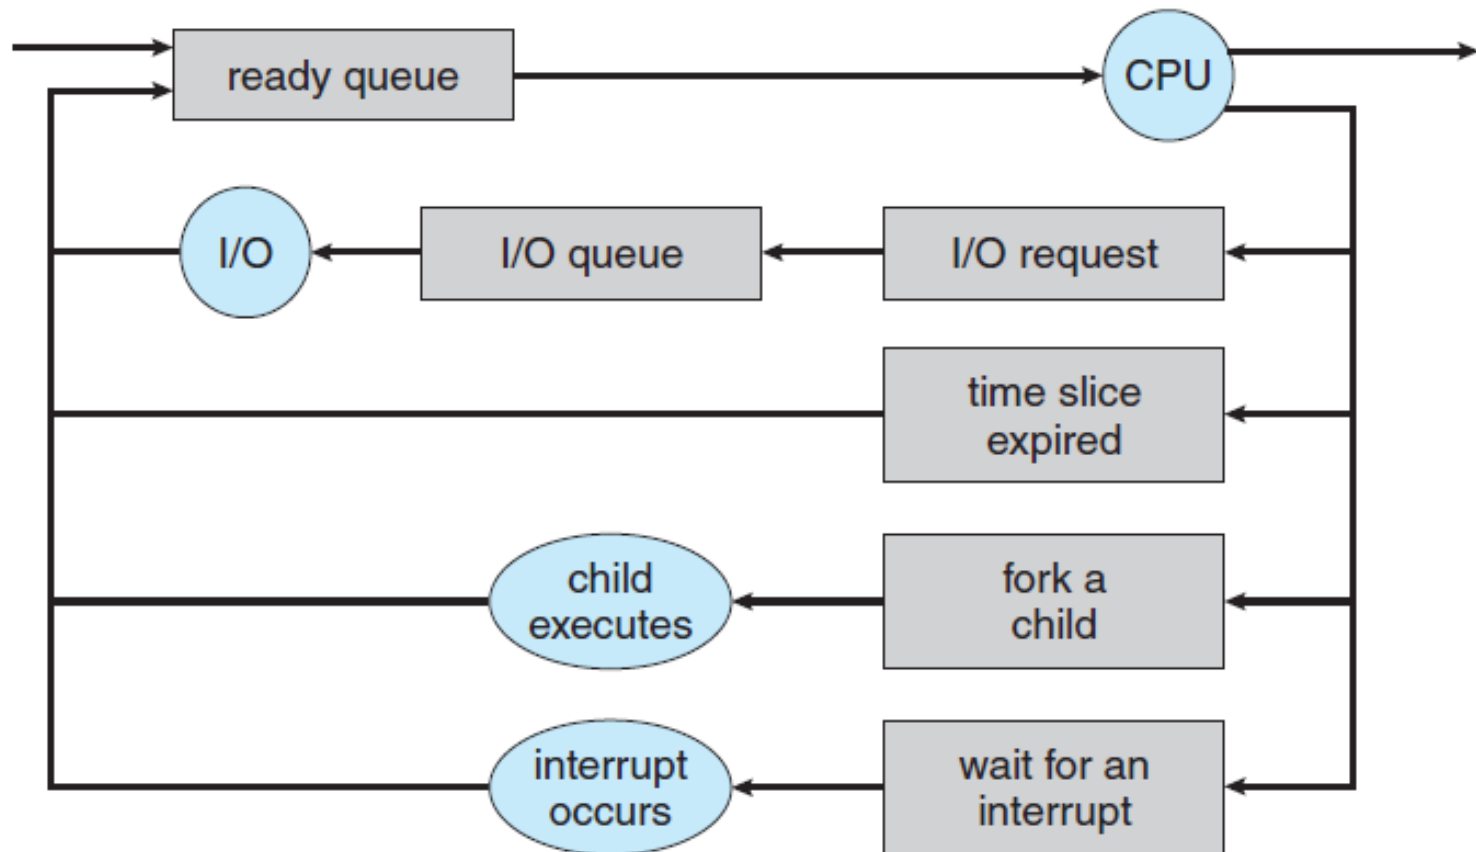



Figure 3.6 Queueing-diagram representation of process scheduling.





Scheduling Queues

- Once the process is allocated to the CPU and is executing, one of several events could occur:
 - The process could issue an **I/O request** and then be placed in an **I/O queue**.
 - The process could create a **new child process** and wait for the child's termination.
 - The process could be **removed forcibly** from the CPU, as a result of an **interrupt**, and be put back in the **ready queue**.
- In the first two cases, the process eventually switches from the **waiting state** to the **ready state** and is then put back in the **ready queue**.
- A process continues this cycle until it terminates.





Schedulers

- A **process** *migrates among the various scheduling queues throughout its lifetime.*
 - The OS must select, for **scheduling purposes**, processes from these queues in some fashion.
 - The selection process for execution is carried out by the appropriate **scheduler**.





Long-term and Short-term schedulers

- ❑ In a **batch system**, more processes are submitted than can be executed immediately (by a single processor).
 - ❑ These processes are **pooled to a mass-storage device**, where they are kept for later execution.
- ❑ A **long-term scheduler**, or **job scheduler**, selects processes from this pool and loads them into memory for execution.
- ❑ A **short-term scheduler**, or **CPU scheduler**, selects the processes from the **ready queue** that are ready to execute and allocates CPU to one of them.
 - ❑ The primary distinction between these two schedulers lies in frequency of execution.





Short-term schedulers

- The **short-term scheduler** selects a new process from the **ready queue** for the CPU frequently.
 - A process may execute for only a few milliseconds before waiting for an **I/O request**.
 - Often, the short-term scheduler executes at least once every 100 milliseconds.
 - If it takes 10 milliseconds to decide to execute a process for 100 milliseconds (total is 110ms), then $10/(110) = 9\%$ of the CPU is being used (wasted) simply for scheduling the work.






Long-term schedulers

- ❑ The **long-term scheduler** executes much **less frequently**.
- ❑ The long-term scheduler controls the **degree of multiprogramming** (*the number of processes in memory*).
- ❑ Because of the longer interval between executions, the **long-term scheduler** can afford to take more time to decide which process should be selected for execution.





I/O-bound process and CPU-bound process

- In general, most processes can be described as either **I/O bound** or **CPU bound**.
 - An **I/O-bound process** is one that spends most of its time doing **I/O** than it spends doing computations.
 - A **CPU-bound process**, generates less I/O requests and using more of its time doing computations.

```
{
printf("\nEnter the first integer: ");
scanf("%d", &a);
printf("\nEnter the second integer: ");
scanf("%d", &b);
} I/O cycle

c = a+b
d = (a*b)-c
e = a-b
f = d/e
} CPU cycle


printf("\n a+b= %d", c);
printf("\n (a*b)-c = %d", d);
printf("\n a-b = %d", e);
printf("\n d/e = %d", f);
} I/O cycle
```





I/O-bound process and CPU-bound process

- The **long-term scheduler** select a good *process mix* of **I/O-bound** and **CPU-bound** processes:
 - If all processes are **I/O bound**, the **ready queue** will almost always be empty,
 - ▶ and the **short-term scheduler** will have little to do.
 - If all processes are **CPU bound**, the **I/O waiting queue** will almost always be empty,
 - ▶ devices will go unused, and again the system will be unbalanced.
- The system with the best performance will thus have a combination of CPU-bound and I/O-bound processes.





Medium-term scheduler (1)

- Some operating systems, such as *time-sharing systems*, may have intermediate level of scheduling, **medium-term scheduler** is shown in **Figure 3.7**.
 - The key idea behind a **medium-term scheduler** is that **it can remove a running process into a special memory** and thus *reduce the degree of multiprogramming*.
 - Later, the process can be reintroduced into ready queue (main memory), and its execution can be continued where it left off.
 - This scheme is called **swapping**. The process is **swapped out**, and is later **swapped in**, by the **medium-term scheduler**.
- **Swapping** is necessary to be free the memory for the **priority** or **pre-empted** process.





Medium-term scheduler (2)

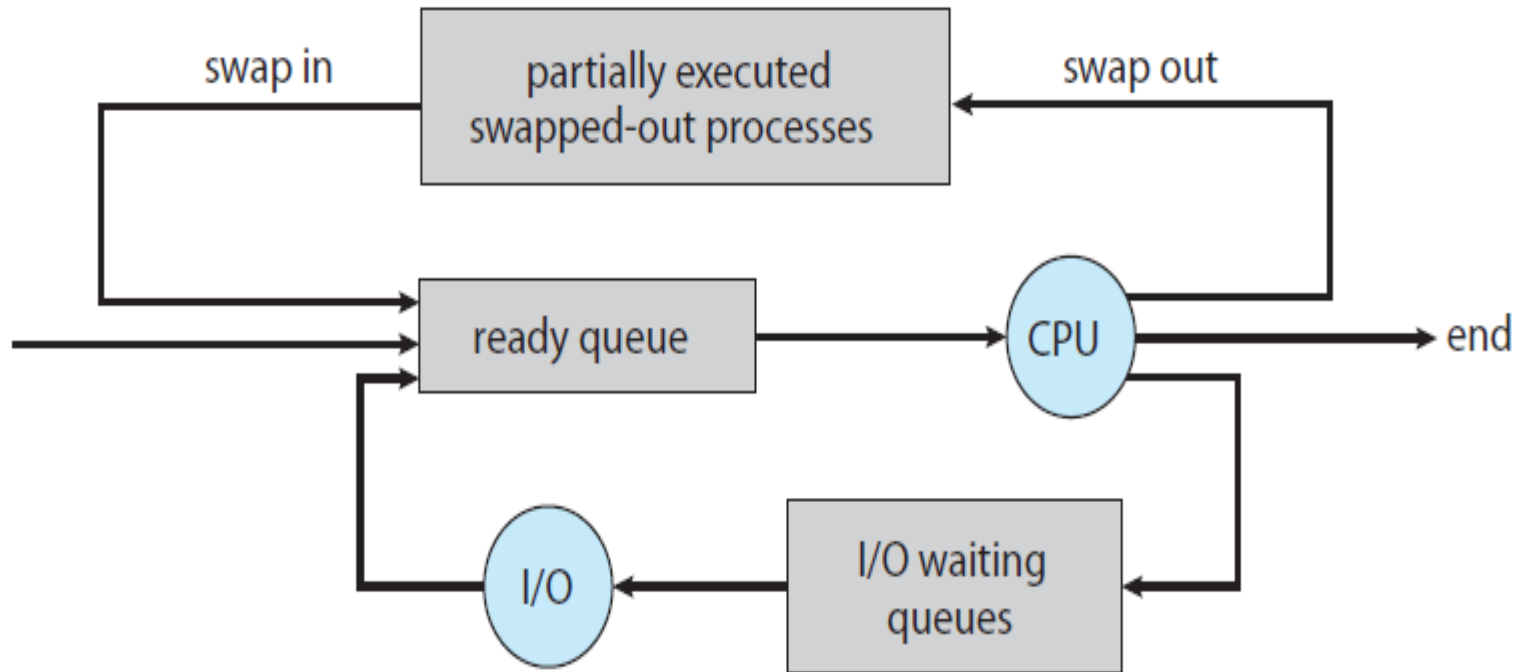


Figure 3.7 Addition of medium-term scheduling to the queueing diagram.





Context Switch (1)

- ❑ When an **interrupt** occurs, the OS needs to save the current **context** of the process running on the CPU so that it can restore that context when its processing is done.
 - ❑ The context is represented in the **PCB** of the process.
 - ❑ It includes the value of the **CPU registers**, the **process state** (see **Figure 3.2**), and **memory-management information**.
- ❑ **Switching** the CPU to another process requires performing a **state save** of the current process and a state restore of a different process. This task is known as a **context switch**.
 - ❑ When a **context switch** occurs, the **kernel** saves the context of the old process in its **PCB** and loads the saved context of the new process scheduled to run.
 - ❑ **Context-switch time** is pure overhead, because the system does no useful work while switching.





Context Switch (3)

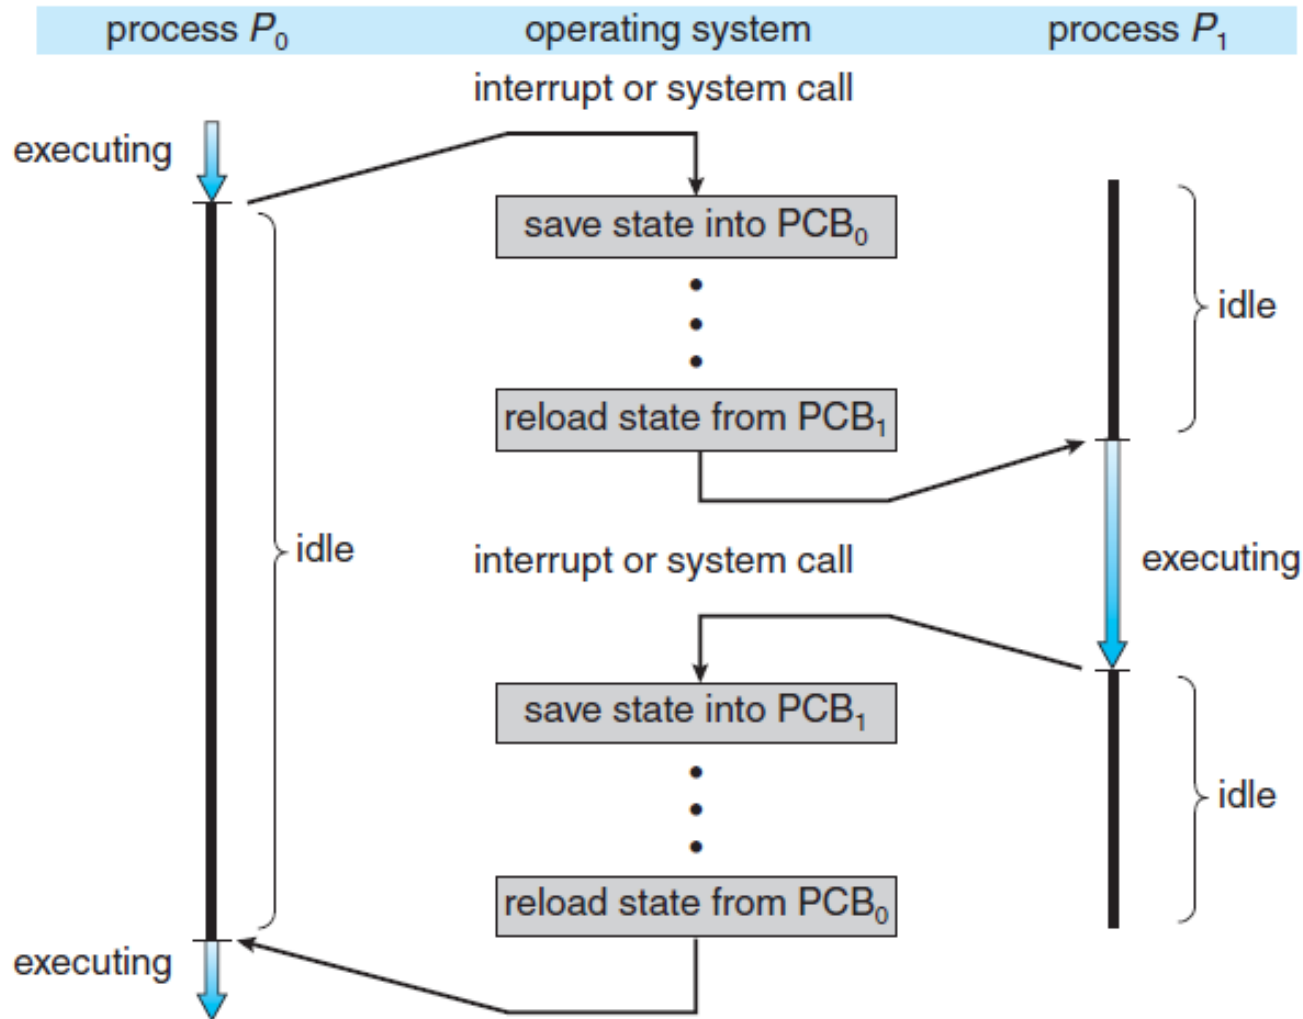


Figure 3.4 Diagram showing CPU switch from process to process.





Process Creation (1)

- During the course of execution, a **process** may create several new processes (called **spawning**).
- The creating process is called a **parent process**, and the new processes are called the **children processes**.
 - Forming a **tree** of processes.
 - Most operating systems including **UNIX**, **Linux**, and **Windows** identify processes according to a unique **process identifier** (or **pid**),
- The **pid** provides a unique value for each process in the system,
 - used as an index to access various attributes of a process within the kernel.





Process Creation (2)

- **Figure 3.8** illustrates a typical process tree for the **Linux** operating system, showing the name of each process and its pid.

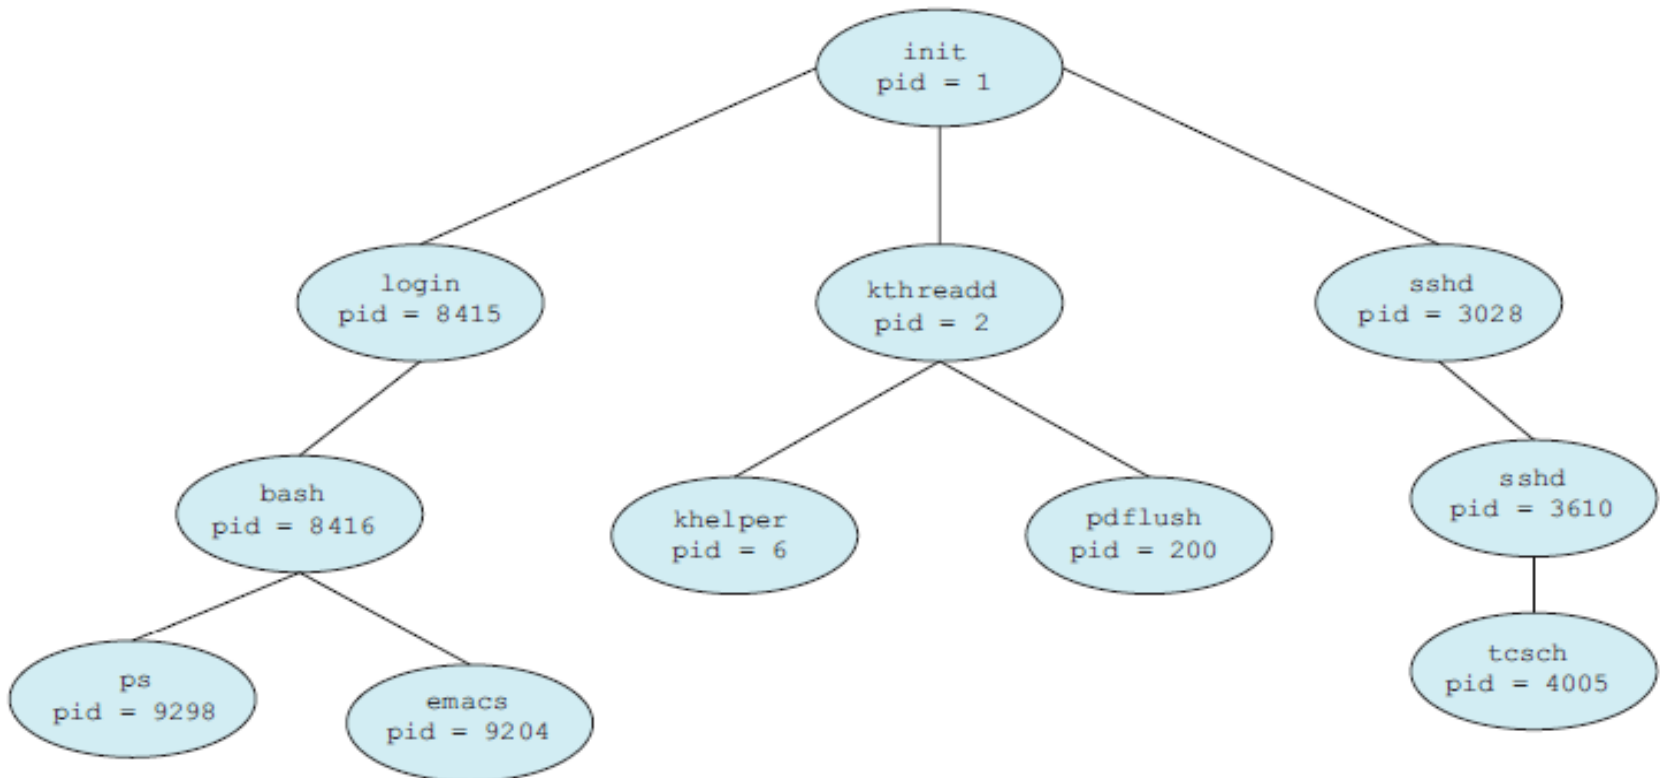


Figure 3.8 A tree of processes on a typical Linux system.





Process Creation (3)

- ❑ The **init** process (which always has a pid of 1) serves as the root parent process for all user processes.
- ❑ Once the system has booted, the **init** process can also create various user processes, such as a web or print server, an **ssh** server, and the like.
- ❑ The **kthreadd** process is responsible for creating additional processes that perform tasks on behalf of the kernel (in this situation, **khelper** and **pdflush**).
- ❑ The **sshd** process is responsible for managing clients that connect to the system by using **ssh** (which is short for **secure shell**).
- ❑ The **login** process is responsible for managing clients that directly log onto the system.
- ❑ In this example, a client has logged on and is using the **bash** shell, which has been assigned pid 8416.
- ❑ Using the **bash** command-line interface, this user has created the process **ps** as well as the **emacs** editor.
- ❑ On UNIX and Linux systems, we can obtain a listing of processes by using the **ps** command. For example, the command **ps -el**






Process Termination

- Process executes last statement and then asks the OS to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are de-allocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates.
 - A process can cause the termination of another process via an appropriate system call (for example **TerminateProcess()** in **Windows**).





Independent and Cooperating processes

- Processes executing concurrently in the OS may be either ***independent processes*** or ***cooperating processes***.
 - A process is ***independent*** if it cannot affect or be affected by the other processes executing in the system.
 - ▶ Any process that does not share data with any other process is independent.
 - A process is ***cooperating*** if it can affect or be affected by the other processes executing in the system.
 - ▶ Clearly, any process that shares data with other processes is a ***cooperating process***.





Process Cooperation

- There are several reasons for providing an environment that allows **process cooperation**:
 - **Information sharing**: several users may be interested in the same piece of information (for instance, a shared file), system must allow concurrent access to such information.
 - **Computation speedup**: If we want a particular task to run faster, we must break it into subtasks and allow them to execute in parallel.
 - **Modularity**: We may want to construct the system in a modular fashion, dividing the system functions into separate processes or threads.
 - **Convenience**: Even an individual user may work on many tasks at the same time. For instance, a user may be editing, listening to music, and compiling in parallel.





Inter-process Communication (IPC)

- ❑ **Cooperating processes** require an **inter-process communication (IPC) mechanism** that will allow them to exchange data and information.
- ❑ There are two fundamental models of IPC:
 - ❑ **shared memory**: In the shared-memory model, a region of memory that is shared by cooperating processes is established.
 - ▶ Processes can then exchange information by reading and writing data to the shared region.
 - ❑ **Message passing**: In the message-passing model, communication takes place by means of messages exchanged between the cooperating processes.
- ❑ The two communications models are contrasted in **Figure 3.12**.





Interprocess Communication

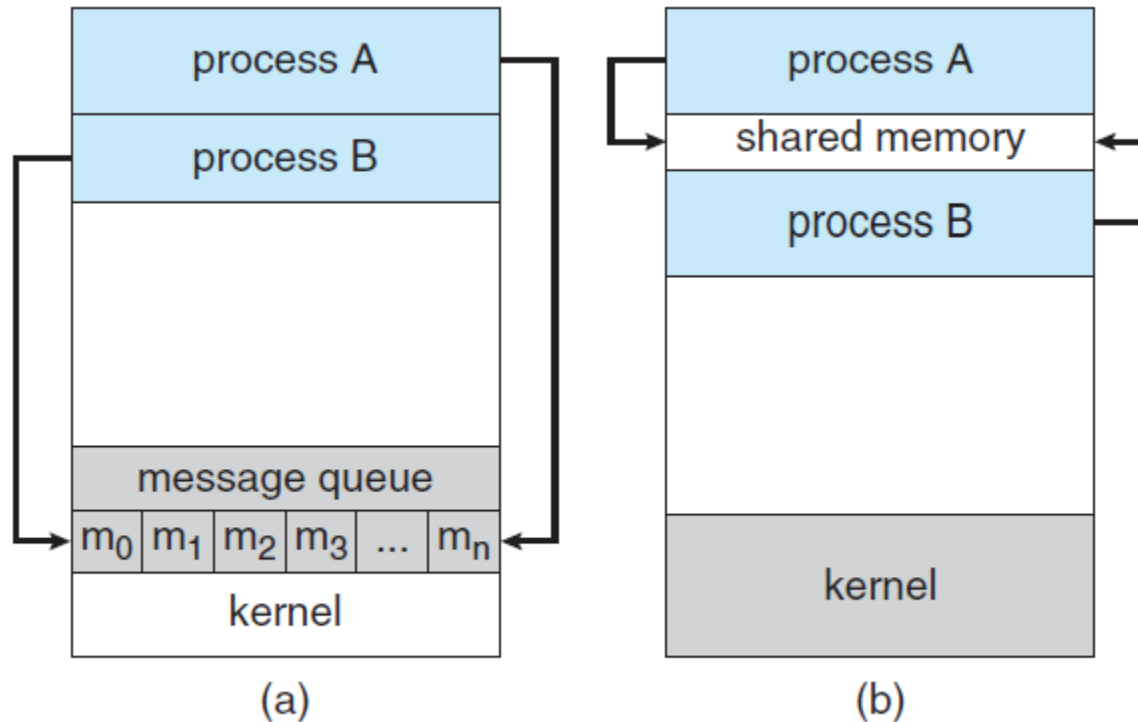


Figure 3.12 Communications models. (a) Message passing. (b) Shared memory.





Interprocess Communication

- ❑ **Message passing** is useful for exchanging smaller amounts of data, because no conflicts need be avoided.
 - ❑ It is easier to implement in a **distributed system** than shared memory.
 - ❑ It needs **system calls** to implement .
 - ❑ Hence it requires more time-consuming task of kernel intervention.
- ❑ **Shared memory** can be faster than message passing, since its does not need system calls to implement.
 - ❑ Here **system calls** are required only to establish **shared memory regions**.
 - ▶ Once shared memory is established, all accesses are treated as routine memory accesses, and no assistance from the kernel is required.





Shared-Memory Systems

- ❑ The OS tries to prevent one process from accessing another process's memory.
 - ❑ **Shared memory** requires that two or more processes agree to remove this restriction.
 - ❑ They can then exchange information by reading and writing data in the shared areas.
 - ❑ The form of the data and the location are determined by these processes and are not under the operating system's control.
 - ❑ The processes are also responsible for ensuring that they are not writing to the same location simultaneously.





Shared-Memory Systems

- To illustrate the concept of **cooperating processes**, let's consider the **producer–consumer** problem, which is a common paradigm for cooperating processes.
 - A **producer process** produces information that is consumed by a consumer process.
 - ▶ For example, a **compiler** may produce **assembly code** that is consumed by an **assembler** section of the compiler.
 - ▶ The assembler, in turn, may produce **object modules** that are consumed by the **loader** section of the assembler.
- The **producer–consumer problem** also provides a useful metaphor for the **client–server** paradigm.





producer–consumer problem

- One solution to the **producer–consumer** problem uses **shared memory**.
 - To allow producer and consumer processes to run concurrently, we must have available a **buffer** of items that can be filled by the producer and emptied by the consumer.
 - This **buffer** will reside in a region of memory that is **shared** by the **producer** and **consumer** processes.
 - A **producer** can produce one item while the **consumer** is consuming another item.
 - The **producer** and **consumer** must be synchronized, so that the consumer does not try to consume an item that has not yet been produced.





producer–consumer problem

- Two types of buffers can be used:
 - **unbounded-buffer** places no practical limit on the size of the buffer
 - **bounded-buffer** assumes that there is a fixed buffer size
- In this case, the consumer must wait if the buffer is empty, and the producer must wait if the buffer is full.
- The code for the producer process is shown in **Figure 3.13**, and the code for the consumer process is shown in **Figure 3.14**.





producer–consumer problem

```
item next_produced;

while (true) {
    /* produce an item in next_produced */

    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Figure 3.13 The producer process using shared memory.





producer–consumer problem

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */
}
```

Figure 3.14 The consumer process using shared memory.





Message Passing

- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
 - ❑ **send**(*message*)
 - ❑ **receive**(*message*)
- ❑ The *message* size is either fixed or variable





Message Passing (Cont.)

- If processes ***P*** and ***Q*** wish to communicate, they need to:
 - Establish a ***communication link*** between them
 - Exchange messages via **send/receive**
- **Implementation issues:**
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?





Message Passing (Cont.)

□ Implementation of communication link

□ **Physical:**

- ▶ Shared memory
- ▶ Hardware bus
- ▶ Network

□ **Logical:**

- ▶ Direct or indirect
- ▶ Synchronous or asynchronous
- ▶ Automatic or explicit buffering





Direct Communication

- Processes must name each other explicitly:
 - **send (P, message)** – send a message to process P
 - **receive(Q, message)** – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication - Mailbox

- ❑ Messages are directed and received from **mailboxes** (also referred to as **ports**)
 - ❑ Each mailbox has a unique **id**
 - ❑ Processes can communicate only if they share the common mailbox
- ❑ Properties of communication link
 - ❑ Link established only if processes share a common mailbox
 - ❑ A link may be associated with many processes
 - ❑ Each pair of processes may share several communication links
 - ❑ Link may be unidirectional or bi-directional





Indirect Communication - Mailbox

□ Operations

- **create** a new mailbox (port)
- **send** and **receive** messages through mailbox
- **destroy** a mailbox

□ Primitives are defined as:

send(*A, message*) – send a message to mailbox *A*

receive(*A, message*) – receive a message from mailbox *A*





Indirect Communication

□ Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 , sends; P_2 and P_3 receive
- Who gets the message?

□ Solutions


- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver.
 - ▶ Sender is notified who the receiver was.





Synchronization

- ❑ Communication between processes takes place through calls to **send()** and **receive()** primitives.
- ❑ There are different design options for implementing each primitive. Message passing may be either **blocking** or **non-blocking**;
- ❑ **Blocking** is considered **synchronous**
 - ❑ **Blocking send** -- the sender is blocked until the message is received
 - ❑ **Blocking receive** -- the receiver is blocked until a message is available
- ❑ **Non-blocking** is considered **asynchronous**
 - ❑ **Non-blocking send** -- the sender sends the message and continue
 - ❑ **Non-blocking receive** -- the receiver receives:
 - ❑ A valid message, or Null message
- ❑ Different combinations possible
 - ❑ If both send and receive are blocking, we have a **rendezvous**






Synchronization (Cont.)

- **With message passing, Producer-consumer becomes trivial**

```
message next_produced;

while (true) {
    /* produce an item in next produced */
    send(next_produced);
    message next_consumed;
} while (true) {
    receive(next_consumed);

    /* consume the item in next consumed */
}
```





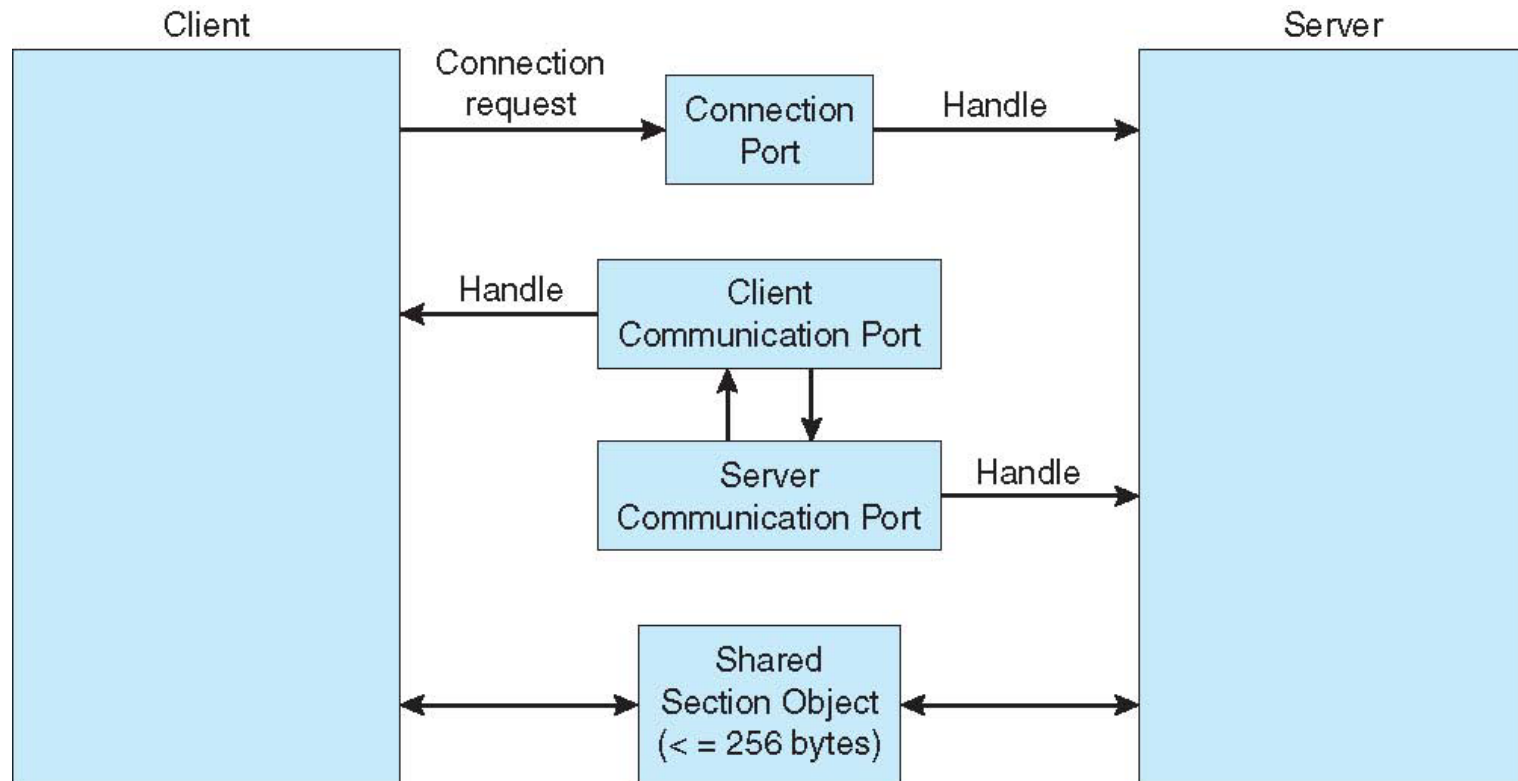
Examples of IPC Systems – Windows

- Message-passing centric via advanced **local procedure call (LPC)** facility
 - ***Only works between processes on the same system***
 - Uses ports (like **mailboxes**) to establish and maintain communication channels
 - Communication works as follows:
 - ▶ The client opens a handle to the subsystem's **connection port** object.
 - ▶ The client sends a **connection request**.
 - ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
 - ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.





Local Procedure Calls in Windows





Remote Procedure Calls

- ❑ **Remote procedure call (RPC)** abstracts *procedure calls* between processes **on networked systems**
 - ❑ Again uses ports for service differentiation
- ❑ **Stubs** – are the intermediate data translation structures
- ❑ The **client-side stub** locates the server and **marshals** (bring together) the parameters
- ❑ The **server-side stub** receives this message, unpacks the marshaled parameters, and performs the procedure on the server
- ❑ On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**





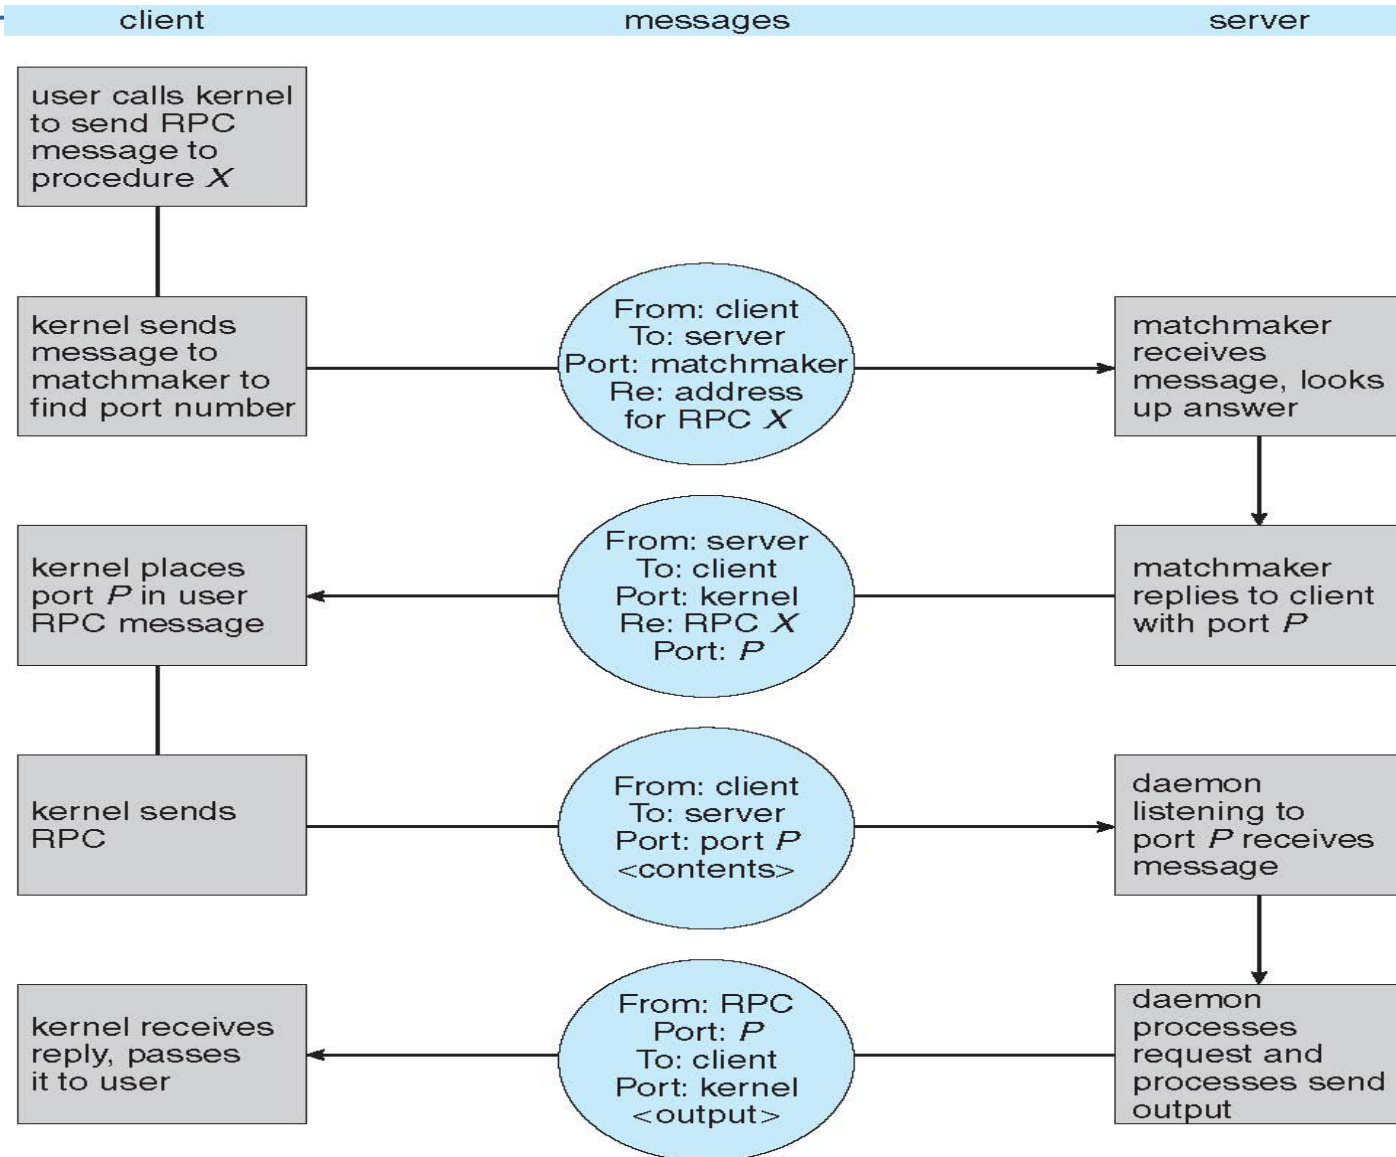
Remote Procedure Calls (Cont.)

- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
 - **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
 - Messages can be delivered ***exactly once*** rather than ***at most once***
- OS typically provides a **rendezvous** (or **matchmaker**) service to connect client and server





Execution of RPC





Programming Exercises

Show **thread creation** in Java using *Executor* and *Runnable* objects.



End of Chapter 3

